

# NY Shooting Incident Data Assignment

CS

## Assignment

Import, tidy and analyze the NYPD Shooting Incident dataset obtained. Be sure your project is reproducible and contains some visualization and analysis. You may use the data to do any analysis that is of interest to you. You should include at least two visualizations and one model. Be sure to identify any bias possible in the data and in your analysis.

```
library(tidyverse)
library(lubridate)
library(forecast)
```

```
-- Attaching core tidyverse packages ----- tidyverse 2.0.0 --
v dplyr      1.1.4      v readr      2.1.5
v forcats    1.0.0      v stringr    1.5.1
v ggplot2    3.5.1      v tibble     3.2.1
v lubridate  1.9.3      v tidyr      1.3.1
v purrr      1.0.2
```

```
-- Conflicts ----- tidyverse_conflicts() --
```

```
x dplyr::filter() masks stats::filter()
x dplyr::lag()     masks stats::lag()
```

```
i Use the conflicted package (<http://conflicted.r-lib.org/>) to force all conflicts to become
Registered S3 method overwritten by 'quantmod':
```

```
  method      from
as.zoo.data.frame zoo
```

```
# import the source data and put it in a df
```

```
source_url <- paste0(
  "https://data.cityofnewyork.us/api/views/833y-fsy8/rows.csv?",
  "accessType=DOWNLOAD"
)
```

```
incident_df <- read.csv(source_url)
```

---

## Explore

```
# see what columns we have and what data types
glimpse(incident_df)
```

```
Rows: 28,562
Columns: 21
$ INCIDENT_KEY      <int> 244608249, 247542571, 84967535, 202853370, 270~
$ OCCUR_DATE        <chr> "05/05/2022", "07/04/2022", "05/27/2012", "09/~
$ OCCUR_TIME        <chr> "00:10:00", "22:20:00", "19:35:00", "21:00:00"~
$ BORO              <chr> "MANHATTAN", "BRONX", "QUEENS", "BRONX", "BROO~
$ LOC_OF_OCCUR_DESC <chr> "INSIDE", "OUTSIDE", "", "", "", "", "", "", ""~
$ PRECINCT          <int> 14, 48, 103, 42, 83, 23, 113, 77, 48, 49, 73, ~
$ JURISDICTION_CODE <int> 0, 0, 0, 0, 0, 2, 0, 0, 0, 0, 0, 0, 0, 0, 0~
$ LOC_CLASSFCTN_DESC <chr> "COMMERCIAL", "STREET", "", "", "", "", "", "", ""~
$ LOCATION_DESC     <chr> "VIDEO STORE", "(null)", "", "", "", "MULTI DW~
$ STATISTICAL_MURDER_FLAG <chr> "true", "true", "false", "false", "false", "fa~
$ PERP_AGE_GROUP    <chr> "25-44", "(null)", "", "25-44", "25-44", "", ""~
$ PERP_SEX          <chr> "M", "(null)", "", "M", "M", "", "", "", "", ""~
$ PERP_RACE          <chr> "BLACK", "(null)", "", "UNKNOWN", "BLACK", "", ~
$ VIC_AGE_GROUP      <chr> "25-44", "18-24", "18-24", "25-44", "25-44", ""~
$ VIC_SEX            <chr> "M", "M", "M", "M", "M", "M", "M", "M", "M", ""~
$ VIC_RACE           <chr> "BLACK", "BLACK", "BLACK", "BLACK", "BLACK", ""~
$ X_COORD_CD         <dbl> 986050, 1016802, 1048632, 1014493, 1009149, 99~
$ Y_COORD_CD         <dbl> 214231.0, 250581.0, 198262.0, 242565.0, 190104~
$ Latitude           <dbl> 40.75469, 40.85440, 40.71063, 40.83242, 40.688~
$ Longitude          <dbl> -73.99350, -73.88233, -73.76777, -73.89071, -7~
$ Lon_Lat            <chr> "POINT (-73.9935 40.754692)", "POINT (-73.8823~
```

I already see lots of nulls, empty strings, missing values, etc. Let's take a closer look at some of the categorical columns to see if there are a limited number of consistently entered values or if they were entered as free text, which might be too difficult to clean.

```
# break out all the unique values and counts
desc_counts <- lapply(incident_df[,
  c(
    "LOC_CLASSFCTN_DESC",
    "LOCATION_DESC",
```

```

"PERP_RACE",
"VIC_RACE",
"LOC_OF_OCCUR_DESC",
"VIC_SEX",
"PERP_SEX"
)
], table)

print(desc_counts)

```

#### \$LOC\_CLASSFCTN\_DESC

	(null)	COMMERCIAL	DWELLING	HOUSING	OTHER
25596	2	208	243	460	59
PARKING LOT	PLAYGROUND	STREET	TRANSIT	VEHICLE	
15	41	1886	23	29	

#### \$LOCATION\_DESC

	(null)	ATM
14977	1711	1
BANK	BAR/NIGHT CLUB	BEAUTY/NAIL SALON
3	668	119
CANDY STORE	CHAIN STORE	CHECK CASH
7	7	1
CLOTHING BOUTIQUE	COMMERCIAL BLDG	DEPT STORE
14	304	9
DOCTOR/DENTIST	DRUG STORE	DRY CLEANER/LAUNDRY
1	14	32
FACTORY/WAREHOUSE	FAST FOOD	GAS STATION
8	130	74
GROCERY/BODEGA	GYM/FITNESS FACILITY	HOSPITAL
750	4	77
HOTEL/MOTEL	JEWELRY STORE	LIQUOR STORE
35	14	42
LOAN COMPANY	MULTI DWELL - APT BUILD	MULTI DWELL - PUBLIC HOUS
1	2964	5007
NONE	PHOTO/COPY STORE	PVT HOUSE
175	1	983
RESTAURANT/DINER	SCHOOL	SHOE STORE
212	1	10
SMALL MERCHANT	SOCIAL CLUB/POLICY LOCATI	STORAGE FACILITY

44	73	1
STORE UNCLASSIFIED	SUPERMARKET	TELECOMM. STORE
37	21	11
VARIETY STORE	VIDEO STORE	
11	8	

\$PERP\_RACE

	(null)
9310	1141
AMERICAN INDIAN/ALASKAN NATIVE	ASIAN / PACIFIC ISLANDER
2	169
BLACK	BLACK HISPANIC
11903	1392
UNKNOWN	WHITE
1837	298
WHITE HISPANIC	
2510	

\$VIC\_RACE

AMERICAN INDIAN/ALASKAN NATIVE	ASIAN / PACIFIC ISLANDER
11	440
BLACK	BLACK HISPANIC
20235	2795
UNKNOWN	WHITE
70	728
WHITE HISPANIC	
4283	

\$LOC\_OF\_OCCUR\_DESC

	INSIDE	OUTSIDE
25596	460	2506

\$VIC\_SEX

F	M	U
2760	25790	12

\$PERP\_SEX

(null)	F	M	U
--------	---	---	---

9310 1141 444 16168 1499

Everything seems to be consistently entered (no misspellings or variations.) But there is a weird mix of “unknown”, “U”, and “null”. It will probably be best to recode empty values as “Unknown” for consistency. There is something weird in a few columns too.

```
unique(incident_df$PERP_RACE)
table(incident_df$PERP_RACE)
```

1. 'BLACK'
2. '(null)'
3. ''
4. 'UNKNOWN'
5. 'WHITE HISPANIC'
6. 'BLACK HISPANIC'
7. 'ASIAN / PACIFIC ISLANDER'
8. 'WHITE'
9. 'AMERICAN INDIAN/ALASKAN NATIVE'

		(null)
	9310	1141
AMERICAN INDIAN/ALASKAN NATIVE		ASIAN / PACIFIC ISLANDER
	2	169
BLACK		BLACK HISPANIC
	11903	1392
UNKNOWN		WHITE
	1837	298
WHITE HISPANIC		
	2510	

Oh, that’s annoying - there is an empty string '' as one of the largest groups, I guess the best option will be to categorize that as "UNKNOWN" as well. While I’m at it I’m going to make the date time columns a little more usable by separating out the date and time and converting them to the right type.

---

## Cleanup

```

clean_incident_df <- incident_df %>%
  mutate(
    # Combine the date and time into a proper DateTime object
    Date = as.POSIXct(
      paste(OCCUR_DATE, OCCUR_TIME),
      format = "%m/%d/%Y %H:%M:%S"
    )
  ) %>%
  rename(
    In_Out = LOC_OF_OCCUR_DESC,
    Location_Category = LOC_CLASSFCTN_DESC,
    Location_details = LOCATION_DESC
  ) %>%
  select(
    Date, BORO, Location_Category, Location_details,
    In_Out, OCCUR_DATE, OCCUR_TIME,
    -JURISDICTION_CODE, -X_COORD_CD, -Y_COORD_CD,
    -Latitude, -Longitude, -Lon_Lat, -PRECINCT,
    everything()
  ) %>%
  mutate(
    # Recode specific values in PERP_RACE and VIC_RACE
    PERP_RACE = recode(PERP_RACE,
      "ASIAN / PACIFIC ISLANDER" = "ASIAN_PAC_ISLAND",
      "AMERICAN INDIAN/ALASKAN NATIVE" = "AM_INDIAN/ALASKAN"
    ),
    VIC_RACE = recode(VIC_RACE,
      "ASIAN / PACIFIC ISLANDER" = "ASIAN_PAC_ISLAND",
      "AMERICAN INDIAN/ALASKAN NATIVE" = "AM_INDIAN/ALASKAN"
    ),
    # Recode empty or null values to "UNKNOWN" for PERP_RACE
    PERP_RACE = ifelse(PERP_RACE == "", "UNKNOWN", PERP_RACE),

    # Recode unknown, empty, or NA values to "U" for both victim and perpetrator sex
    PERP_SEX = ifelse(PERP_SEX %in% c("Unknown", "", "(null)", NA), "U", PERP_SEX),
    VIC_SEX = ifelse(VIC_SEX %in% c("Unknown", "", NA), "U", VIC_SEX)
  )

# Check that the recoding worked as expected
unique(clean_incident_df$PERP_SEX) # Should show only "M", "F", and "U"
unique(clean_incident_df$VIC_SEX)  # Should show only "M", "F", and "U"

```

```
# check that I have the columns and order that I wanted
glimpse(clean_incident_df)

# check that we fixed the empty string values
unique(clean_incident_df$PERP_RACE)

# printing a df is a little uglier in some ways but prevents text
# overlap when there are lots of columns or long column names
print(tail(clean_incident_df))
```

1. 'M'
2. 'U'
3. 'F'

1. 'M'
2. 'F'
3. 'U'

Rows: 28,562

Columns: 22

```
$ Date          <dtm> 2022-05-05 00:10:00, 2022-07-04 22:20:00, 201~
$ BORO          <chr> "MANHATTAN", "BRONX", "QUEENS", "BRONX", "BROO~
$ Location_Category <chr> "COMMERCIAL", "STREET", "", "", "", "", "", ""~
$ Location_details <chr> "VIDEO STORE", "(null)", "", "", "", "MULTI DW~
$ In_Out        <chr> "INSIDE", "OUTSIDE", "", "", "", "", "", "", ""~
$ OCCUR_DATE    <chr> "05/05/2022", "07/04/2022", "05/27/2012", "09/~
$ OCCUR_TIME    <chr> "00:10:00", "22:20:00", "19:35:00", "21:00:00"~
$ INCIDENT_KEY  <int> 244608249, 247542571, 84967535, 202853370, 270~
$ PRECINCT      <int> 14, 48, 103, 42, 83, 23, 113, 77, 48, 49, 73, ~
$ JURISDICTION_CODE <int> 0, 0, 0, 0, 0, 2, 0, 0, 0, 0, 0, 0, 0, 0, 0~
$ STATISTICAL_MURDER_FLAG <chr> "true", "true", "false", "false", "false", "fa~
$ PERP_AGE_GROUP <chr> "25-44", "(null)", "", "25-44", "25-44", "", "~
$ PERP_SEX      <chr> "M", "U", "U", "M", "M", "U", "U", "U", "U", "~
$ PERP_RACE     <chr> "BLACK", "(null)", "UNKNOWN", "UNKNOWN", "BLAC~
$ VIC_AGE_GROUP <chr> "25-44", "18-24", "18-24", "25-44", "25-44", "~
$ VIC_SEX       <chr> "M", "M", "M", "M", "M", "M", "M", "M", "M", "~
$ VIC_RACE      <chr> "BLACK", "BLACK", "BLACK", "BLACK", "BLACK", "~
$ X_COORD_CD    <dbl> 986050, 1016802, 1048632, 1014493, 1009149, 99~
$ Y_COORD_CD    <dbl> 214231.0, 250581.0, 198262.0, 242565.0, 190104~
$ Latitude      <dbl> 40.75469, 40.85440, 40.71063, 40.83242, 40.688~
$ Longitude     <dbl> -73.99350, -73.88233, -73.76777, -73.89071, -7~
$ Lon_Lat       <chr> "POINT (-73.9935 40.754692)", "POINT (-73.8823~
```

1. 'BLACK'
2. '(null)'
3. 'UNKNOWN'
4. 'WHITE HISPANIC'
5. 'BLACK HISPANIC'
6. 'ASIAN\_PAC\_ISLAND'
7. 'WHITE'
8. 'AM\_INDIAN/ALASKAN'

	Date	BORO	Location_Category	Location_details		
28557	2023-07-02 21:40:00	BRONX	STREET	(null)		
28558	2023-03-19 23:48:00	BRONX	COMMERCIAL	GROCERY/BODEGA		
28559	2023-08-16 02:46:00	BRONX	STREET	(null)		
28560	2023-06-27 12:27:00	BRONX	DWELLING	MULTI DWELL - APT BUILD		
28561	2023-07-08 11:27:00	QUEENS	STREET	BEAUTY/NAIL SALON		
28562	2023-07-24 23:38:00	MANHATTAN	HOUSING	MULTI DWELL - PUBLIC HOUS		
	In_Out	OCCUR_DATE	OCCUR_TIME	INCIDENT_KEY	PRECINCT	JURISDICTION_CODE
28557	OUTSIDE	07/02/2023	21:40:00	270719378	46	0
28558	INSIDE	03/19/2023	23:48:00	265354835	47	0
28559	OUTSIDE	08/16/2023	02:46:00	272968931	41	0
28560	INSIDE	06/27/2023	12:27:00	270489846	41	0
28561	OUTSIDE	07/08/2023	11:27:00	271021661	102	0
28562	OUTSIDE	07/24/2023	23:38:00	271818283	28	2
	STATISTICAL_MURDER_FLAG	PERP_AGE_GROUP	PERP_SEX	PERP_RACE		
28557	false	(null)	U	(null)		
28558	true	18-24	M	BLACK		
28559	false	25-44	F	BLACK		
28560	true	25-44	M	BLACK		
28561	false	25-44	M	WHITE HISPANIC		
28562	false	(null)	U	(null)		
	VIC_AGE_GROUP	VIC_SEX	VIC_RACE	X_COORD_CD	Y_COORD_CD	Latitude
28557	18-24	M	BLACK HISPANIC	1009601	247515	40.84601
28558	18-24	M	BLACK	1025687	268586	40.90378
28559	45-64	M	BLACK	1014639	240066	40.82555
28560	25-44	M	BLACK	1012221	238552	40.82140
28561	65+	M	ASIAN_PAC_ISLAND	1028856	192785	40.69572
28562	25-44	M	BLACK	997853	230889	40.80040
	Longitude	Lon_Lat				
28557	-73.90837 POINT (-73.908369 40.846012)					
28558	-73.85010 POINT (-73.850098 40.903785)					
28559	-73.89020 POINT (-73.890195 40.825549)					
28560	-73.89894 POINT (-73.898938 40.821404)					



```
28561 -73.83914 POINT (-73.839138 40.695717)
28562 -73.95086 POINT (-73.950864 40.800405)
```

I'm going to make a few different dataframes with different groups for eventual analysis and plotting. Things I'm going to start with

- Daily incidents over time to look for general trends
- Incidents by month and year
- Incidents by borough
- Incidents by month (not over time, so total incidents that occurred in each month summed over all years)
- Time and year data broken down by borough
- Victim and perpetrator by sex

```
# for plotting incidents over time
time_series_df <- clean_incident_df %>%
  mutate(simple_date = as.Date(OCCUR_DATE, format = "%m/%d/%Y")) %>%
  group_by(simple_date) %>%
# Add a new column that represents only the month and year
# This step may be unnecessary since I have a good date column
# but it's easier for me to understand
  summarise(total_by_day = n()) %>%
  mutate(month_year = floor_date(simple_date, "month"))

# for plotting overtime by month and year
df_aggregated <- time_series_df %>%
  mutate(year = format(simple_date, "%Y"),
         month = format(simple_date, "%m")) %>%
  group_by(year, month) %>%
  summarise(total_by_month = sum(total_by_day)) %>%
  mutate(
    year = as.numeric(year), # Convert year to numeric
    month = as.numeric(month) # Convert month to numeric
  ) %>%
  ungroup()

# borough totals
total_by_borough <- clean_incident_df %>%
  group_by(BORO) %>%
  summarize(total_incidents = n())

# monthly borough totals
```

```

monthly_totals_by_borough <- clean_incident_df %>%
  mutate(month = floor_date(Date, "month")) %>%
  mutate(month = as.Date(month)) %>%
  group_by(BORO, month) %>%
  summarize(monthly_incidents = n()) %>%
  ungroup()

# victim and perp by sex
totals_by_sex <- clean_incident_df %>%
  group_by(PERP_SEX, VIC_SEX) %>%
  summarise(
    Total_Victims = n(),      # Count number of victims in each group
    Total_Perps = n()        # Count number of perpetrators in each group
  )

tail(time_series_df)
tail(df_aggregated)
tail(total_by_borough)
tail(monthly_totals_by_borough)
tail(totals_by_sex, 9)

```

`summarise()` has grouped output by 'year'. You can override using the  
`.groups` argument.

`summarise()` has grouped output by 'BORO'. You can override using the  
`.groups` argument.

`summarise()` has grouped output by 'PERP\_SEX'. You can override using the  
`.groups` argument.

A tibble: 6 x 3

simple_date <date>	total_by_day <int>	month_year <date>
2023-12-22	8	2023-12-01
2023-12-23	4	2023-12-01
2023-12-24	5	2023-12-01
2023-12-26	6	2023-12-01
2023-12-27	1	2023-12-01
2023-12-29	3	2023-12-01

A tibble: 6 x 3

year <dbl>	month <dbl>	total_by_month <int>
2023	7	152
2023	8	108
2023	9	105
2023	10	99
2023	11	71
2023	12	83

A tibble: 5 x 2

BORO <chr>	total_incidents <int>
BRONX	8376
BROOKLYN	11346
MANHATTAN	3762
QUEENS	4271
STATEN ISLAND	807

A tibble: 6 x 3

BORO <chr>	month <date>	monthly_incidents <int>
STATEN ISLAND	2023-05-01	3
STATEN ISLAND	2023-06-01	8
STATEN ISLAND	2023-07-01	6
STATEN ISLAND	2023-08-01	3
STATEN ISLAND	2023-10-01	3
STATEN ISLAND	2023-11-01	2

A grouped\_df: 9 x 4

PERP_SEX <chr>	VIC_SEX <chr>	Total_Victims <int>	Total_Perps <int>
F	F	77	77
F	M	366	366
F	U	1	1
M	F	1755	1755
M	M	14406	14406
M	U	7	7
U	F	928	928

PERP_SEX <chr>	VIC_SEX <chr>	Total_Victims <int>	Total_Perps <int>
U	M	11018	11018
U	U	4	4

That's looks pretty good. I'm going to make all my plots at once, so it will get a little messy looking, but that will be the easiest way to set some universal configurations (theme, size, etc.) Then we can use these to decide on further plotting or analysis or modeling to do.

## Visualization

```
options(repr.plot.width = 10, repr.plot.height = 7)
theme_set(theme_minimal())

# plot daily incidents
ggplot(time_series_df, aes(x = simple_date, y = total_by_day)) +
  geom_line(color = "skyblue") +
  labs(
    title = "Fig.1 - Daily Incidents",
    x = "Date",
    y = "Incident Count"
  )

# plot monthly incidents over time
ggplot(time_series_df, aes(x = month_year, y = total_by_day)) +
  geom_bar(stat = "identity", fill = "skyblue") +
  labs(
    title = "Fig. 2 - Total Occurrences by Month",
    x = "Year",
    y = "Total Occurrences") +
  scale_x_date(date_labels = "%Y", date_breaks = "1 year")

# plot yearly incidents
ggplot(time_series_df, aes(
  x = year(simple_date),
  y = total_by_day)) +
```

```

geom_bar(stat = "identity", fill = "skyblue") +
labs(
  title = "Fig. 3 - Total Occurrences by Year",
  x = "Year",
  y = "Total Occurrences")

# Plot occurrences by month (across all years)
ggplot(time_series_df, aes(
  x = month(simple_date, label = TRUE),
  y = total_by_day
)) +
geom_bar(
  stat = "identity",
  fill = "skyblue"
) +
labs(
  title = "Fig. 4 - Total Occurrences by Month",
  x = "Month",
  y = "Total Occurrences"
)

# Plot total by borough
ggplot(total_by_borough, aes(
  x = BORO,
  y = total_incidents
)) +
geom_bar(
  stat = "identity",
  fill = "skyblue"
) +
labs(
  title = "Fig. 5 - Total by Borough",
  x = "Borough",
  y = "Total"
)

# Borough totals monthly
ggplot(monthly_totals_by_borough, aes(
  x = month,
  y = monthly_incidents,

```

```

color = BORO)
) +
# I wanted to try a line instead of a bar
geom_line(linewidth = 1.2) +
geom_point(size = 2) +
labs(
  title = "Fig. 6 - Monthly Incidents by Borough",
  x = "Month",
  y = "Total Incidents"
) +
scale_x_date(date_labels = "%b %Y", date_breaks = "6 month") +
theme(axis.text.x = element_text(angle = 45, hjust = 1))

# Plot incidents by year for each borough with side-by-side bars
ggplot(monthly_totals_by_borough, aes(
  x = year(month),
  y = monthly_incidents,
  fill = BORO)
) +
# Use dodge for side-by-side bars
geom_bar(stat = "identity", position = "dodge") +
labs(
  title = "Fig. 7 - Yearly Incidents by Borough",
  x = "Year",
  y = "Total Incidents")

# plot for total victims of each sex by perpetrator sex
ggplot(totals_by_sex, aes(
  x = PERP_SEX,
  y = Total_Victims,
  fill = VIC_SEX)
) +
geom_bar(stat = "identity", position = "dodge") +
labs(
  title = "Fig. 8 - Victim Sex by Perpetrator Sex",
  x = "Perpetrator Sex",
  y = "Total Victims",
  fill = "Victim Sex"
) +
theme_minimal()

```

Fig.1 - Daily Incidents




Fig. 2 - Total Occurrences by Month






Fig. 3 - Total Occurrences by Year




Fig. 4 - Total Occurrences by Month




Fig. 5 - Total by Borough




Fig. 6 - Monthly Incidents by Borough



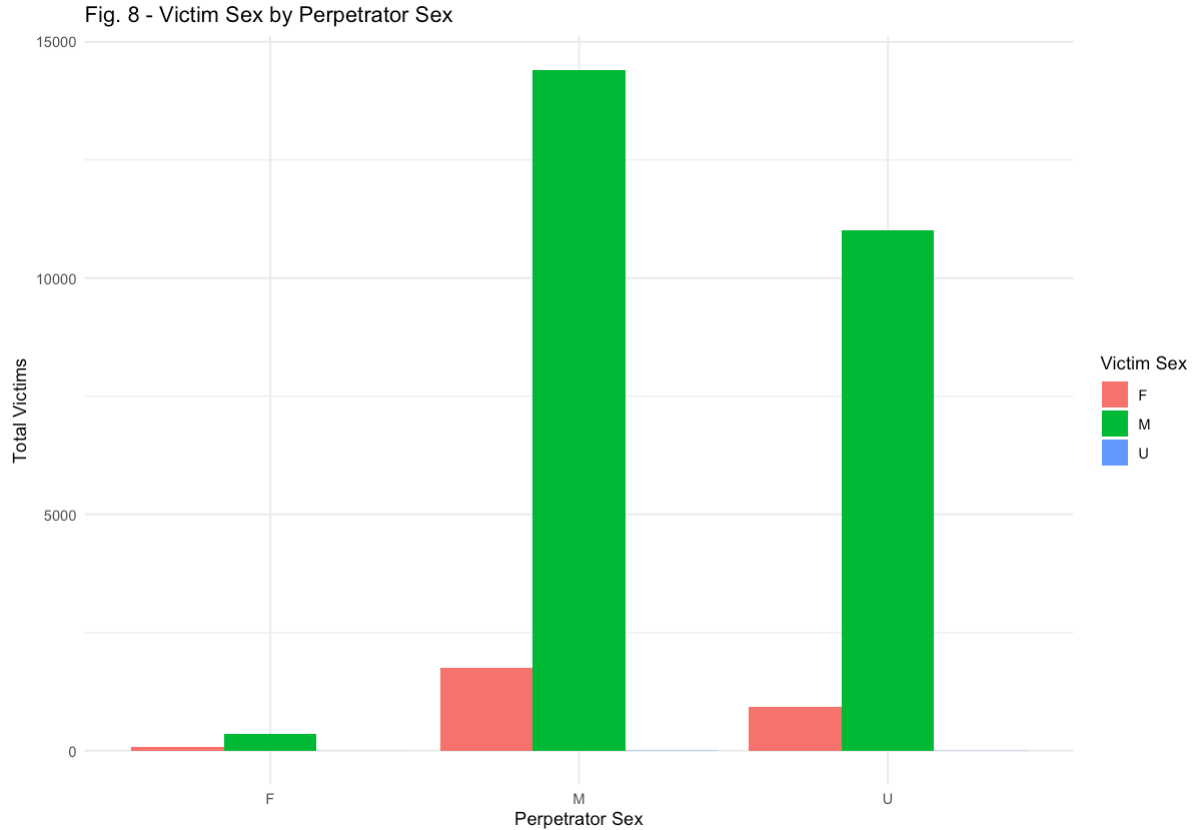


Fig. 7 - Yearly Incidents by Borough





- Fig 1 - Daily incidents - hard to see many trends since the data is so noisy, but it does look like it's generally periodic and there is a big spike around the first covid summer.
- Fig 2 - Total occurrences by month over time - clearer seasonality and a little easier to see the pre- and post-covid trends
- Fig 3 - Total occurrences by year - now we can see trends. Decreasing incidents starting in the early 20-teens and flattening out, before a big covid spike and almost back down to pre-covid levels.
- Fig 4 - Total occurrences by month - pretty strong visual trend towards higher incidents in the hottest months, which is a well studied phenomenon.
- Fig 5 - Total by borough - generally interesting, but would be more useful with the context of per capita and per area data for the boroughs.
- Fig 6/7 - Boroughs over time (final 2 plots) - interesting to see that not all boroughs follow the same trends over time, and that the first covid spike was driven heavily by increased in Brooklyn.
- Fig 8 - Totals of victim sex group by perpetrator sex - there are fewer total female victims than I would have expected.

We could keep going with similar visuals (breakdown by race, gender, age group, etc. or relationships between group like victim age relative to perpetrator age) but I'll stop there.

I'm going to focus on victim sex for the analysis and modeling component. I want to see how predictive of a victim being female some of the other attributes (perpetrator race and sex, victim race). I'll start with a simple model of victim sex as predicted by perpetrator sex. To start will have to exclude the "unknowns" from victim sex and clean up some of the other factors.

---

## Modeling

```
# Convert all categorical variables to factors
clean_incident_df$VIC_SEX <- factor(clean_incident_df$VIC_SEX)
clean_incident_df$BORO <- factor(clean_incident_df$BORO)
clean_incident_df$VIC_RACE <- factor(clean_incident_df$VIC_RACE)
clean_incident_df$PERP_RACE <- factor(clean_incident_df$PERP_RACE)
clean_incident_df$PERP_SEX <- factor(clean_incident_df$PERP_SEX)

# Apply droplevels to all factor columns (to remove unused levels)
clean_incident_df <- clean_incident_df %>%
  mutate_if(is.factor, droplevels)

# Set "M" as the victim reference so that we model the odds of being "Female"
# Set the reference race for victims and perps as "white"
clean_incident_df$VIC_SEX <- relevel(clean_incident_df$VIC_SEX, ref = "M")
clean_incident_df$VIC_RACE <- relevel(clean_incident_df$VIC_RACE, ref = "WHITE")
clean_incident_df$PERP_RACE <- relevel(clean_incident_df$PERP_RACE, ref = "WHITE")

# Check the levels of the factor to confirm they are correct
levels(clean_incident_df$VIC_SEX)
levels(clean_incident_df$PERP_SEX)
levels(clean_incident_df$PERP_RACE)
```

1. 'M'
2. 'F'
3. 'U'

1. 'F'
2. 'M'
3. 'U'

1. 'WHITE'

2. '(null)'
3. 'AM\_INDIAN/ALASKAN'
4. 'ASIAN\_PAC\_ISLAND'
5. 'BLACK'
6. 'BLACK HISPANIC'
7. 'UNKNOWN'
8. 'WHITE HISPANIC'

```
# create a simple generalize linear model to predict odds of female victim
# based on perp sex
simple_model_vic_sex <- glm(VIC_SEX ~ PERP_SEX,
  family=binomial,
  data=clean_incident_df,
  na.action = na.exclude)

summary(simple_model_vic_sex)
```

Call:

```
glm(formula = VIC_SEX ~ PERP_SEX, family = binomial, data = clean_incident_df,
  na.action = na.exclude)
```

Coefficients:

	Estimate	Std. Error	z value	Pr(> z )
(Intercept)	-1.5459	0.1247	-12.396	< 2e-16 ***
PERP_SEXM	-0.5553	0.1272	-4.364	1.28e-05 ***
PERP_SEXU	-0.9240	0.1293	-7.147	8.88e-13 ***

---

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 18197 on 28561 degrees of freedom

Residual deviance: 18093 on 28559 degrees of freedom

AIC: 18099

Number of Fisher Scoring iterations: 5

It looks like each perpetrator sex option has a statistically significant effect on the change in log-odds of the victim being female, so I'll pull out each of them and convert them to odds.



```

print("Log-odds are")
print(coef(simple_model_vic_sex)[c("PERP_SEXM", "PERP_SEXU")])

#convert log-odds to odds and print
print("The odds relative to the female victim/female perpetrator baseline are")
print((exp(coef(simple_model_vic_sex)[c("PERP_SEXM", "PERP_SEXU")]))))

# Print odds for when the perpetrator is female
paste(
  "The odds of a victim being female when the perpetrator is female are",
  round(exp(coef(simple_model_vic_sex)["(Intercept)"]), 2)
)

# Print odds for when the perpetrator is male
paste0(
  "The odds of a victim being female when the perpetrator is male are ",
  round(exp(coef(simple_model_vic_sex)["PERP_SEXM"]), 2) * 100,
  "% of baseline"
)

# Print odds for when the perpetrator is unknown
paste0(
  "The odds of a victim being female when the perpetrator is unknown are ",
  round(exp(coef(simple_model_vic_sex)["PERP_SEXU"]), 2) * 100,
  "% of baseline"
)

```

```

[1] "Log-odds are"
   PERP_SEXM PERP_SEXU
-0.5552708 -0.9240283
[1] "The odds relative to the female victim/female perpetrator baseline are"
   PERP_SEXM PERP_SEXU
0.5739169 0.3969169

```

‘The odds of a victim being female when the perpetrator is female are 0.21’

‘The odds of a victim being female when the perpetrator is male are 57% of baseline’

‘The odds of a victim being female when the perpetrator is unknown are 40% of baseline’

```

# make a better df for monthly modeling
monthly_totals <- time_series_df %>%

```

```

group_by(month_year) %>%
  summarise(total_by_month = sum(total_by_day))

# Check the aggregated data
tail(monthly_totals)

```

A tibble: 6 x 2

month_year <date>	total_by_month <int>
2023-07-01	152
2023-08-01	108
2023-09-01	105
2023-10-01	99
2023-11-01	71
2023-12-01	83


```

# make a simple ratio model based on high and low seasons
# Aggregate by year and calculate sum for January-March and
seasonal_totals <- df_aggregated %>%
  mutate(season = case_when(
    month %in% c(1, 2, 3) ~ "low_season",
    month %in% c(6, 7, 8) ~ "high_season",
    TRUE ~ NA_character_
  )) %>%
  filter(!is.na(season)) %>% # Filter to keep only rows with season
  group_by(year, season) %>%
  summarise(total = sum(total_by_month, na.rm = TRUE), .groups = "drop") %>%
  pivot_wider(names_from = season, values_from = total) %>%
  mutate(ratio = (high_season / low_season) * 100)

# plot the simple ratio results
ggplot(seasonal_totals, aes(
  x = year,
  y = ratio
)) +
  geom_bar(
    stat = "identity",
    fill = "skyblue"
  ) +
  labs(
    title = "High Season (Jun-Aug) vs Low season (Jan-Mar)",

```

```
x = "Year",
y = "% of Low Season"
)
```



```
# make time series friendly
start_year <- as.numeric(format(min(monthly_totals$month_year), "%Y"))
start_month <- as.numeric(format(min(monthly_totals$month_year), "%m"))

# Convert to a time series object (monthly frequency)
monthly_incident_ts <- ts(monthly_totals$total_by_month, start = c(start_year, start_month),

# Check the time series
print(monthly_incident_ts)
```

	Jan	Feb	Mar	Apr	May	Jun	Jul	Aug	Sep	Oct	Nov	Dec
2006	129	97	102	156	173	180	233	245	196	199	167	178
2007	109	78	103	117	199	201	239	212	173	197	135	124

2008	114	80	128	153	194	197	234	215	197	153	130	164
2009	105	113	104	136	172	182	211	210	167	143	161	124
2010	97	123	132	167	166	219	211	203	182	144	140	128
2011	102	98	89	136	176	178	221	217	250	188	133	151
2012	114	101	130	127	173	183	247	191	140	113	98	100
2013	119	66	85	94	99	139	141	162	120	125	98	91
2014	107	52	109	98	129	153	161	169	109	130	107	140
2015	117	68	93	106	181	135	132	162	115	103	104	118
2016	79	74	79	72	108	116	112	184	128	105	77	74
2017	68	48	56	81	88	113	104	96	90	89	61	76
2018	88	41	56	68	98	90	114	92	82	83	71	75
2019	63	59	65	67	81	107	120	105	86	80	71	63
2020	94	51	71	69	137	265	325	311	189	165	139	132
2021	95	89	130	181	227	199	204	228	192	159	151	156
2022	116	102	156	155	170	172	229	154	156	103	101	102
2023	93	104	109	85	111	130	152	108	105	99	71	83

```
sarima_forecast <- auto.arima(monthly_incident_ts, seasonal = TRUE)

# Summary of the SARIMA model
summary(sarima_forecast)

# Forecasting the next 12 months (or any desired period)
forecast_values <- forecast::forecast(sarima_forecast, h = 12)

# Plot the forecast
autoplot(forecast_values) +
  labs(x = "Date", y = "Value", title = "SARIMA Forecast") +
  theme(axis.text.x = element_text(angle = 45, hjust = 1))

# create a zoomed in version
autoplot(forecast_values) +
  labs(title = "Zoomed Forecast", x = "Year", y = "Totals") +
  xlim(c(2019, 2025)) + # Adjust x-axis to show the previous 12 months and forecast
  theme(axis.text.x = element_text(angle = 45, hjust = 1))
```

Series: monthly\_incident\_ts  
 ARIMA(1,0,0)(1,1,0)[12] with drift

Coefficients:

ar1	sar1	drift
0.6957	-0.3703	-0.3007

s.e. 0.0507 0.0662 0.4204

$\sigma^2 = 906.9$ : log likelihood = -983.79

AIC=1975.58 AICc=1975.78 BIC=1988.85

Training set error measures:

	ME	RMSE	MAE	MPE	MAPE	MASE
Training set	-0.01241811	29.05056	21.49203	-2.150574	17.37663	0.7560569

ACF1


Training set -0.07634969


Scale for x is already present.

Adding another scale for x, which will replace the existing scale.

Warning message:

"Removed 156 rows containing missing values or values outside the scale range  
(`geom\_line()`)."





```
# Print the confidence intervals for the next 5 months
next_12_conf_intervals <- data.frame(
  month2024 = c(1:12),
  forecast = forecast_values$mean[1:12],
  lower_80 = forecast_values$lower[1:12, 1], # 80% lower bound
  upper_80 = forecast_values$upper[1:12, 1], # 80% upper bound
  lower_95 = forecast_values$lower[1:12, 2],
  upper_95 = forecast_values$upper[1:12, 2]
) %>%
  mutate(across(everything(), \(x) round(x, 0)))
next_12_conf_intervals
```

A data.frame: 12 x 6

month2024 <dbl>	forecast <dbl>	lower_80 <dbl>	upper_80 <dbl>	lower_95 <dbl>	upper_95 <dbl>
1	73	34	111	14	132
2	82	35	129	10	154
3	110	59	161	33	187
4	98	46	150	18	178
5	122	69	175	41	203
6	137	83	190	55	218
7	173	119	226	91	255
8	118	65	172	36	200
9	118	64	171	36	200
10	95	41	148	12	177
11	77	23	130	-6	159
12	85	31	138	2	167

That's the end of the analysis I'm comfortable with. Below I wanted to see what it would look like to do similar modeling with multiple predictor variables (perp sex, race, borough.) It looks like it worked, but it gets out of hand to interpret it pretty quickly so I just stopped and left it here as an interesting example of what else could be done.

```
model_vic_sex <- glm(VIC_SEX ~ BORO + VIC_RACE +
  PERP_SEX + PERP_RACE,
  family=binomial,
  data=clean_incident_df,
  na.action = na.exclude)

print(levels(clean_incident_df$VIC_RACE))
print(levels(clean_incident_df$PERP_RACE))
print(levels(clean_incident_df$BORO))
print(summary(model_vic_sex))
```

```
[1] "WHITE"           "AM_INDIAN/ALASKAN" "ASIAN_PAC_ISLAND"
[4] "BLACK"           "BLACK HISPANIC"    "UNKNOWN"
[7] "WHITE HISPANIC"
[1] "WHITE"           "(null)"             "AM_INDIAN/ALASKAN"
[4] "ASIAN_PAC_ISLAND" "BLACK"              "BLACK HISPANIC"
[7] "UNKNOWN"         "WHITE HISPANIC"
[1] "BRONX"           "BROOKLYN"          "MANHATTAN"         "QUEENS"
[5] "STATEN ISLAND"
```

Call:

```
glm(formula = VIC_SEX ~ BORO + VIC_RACE + PERP_SEX + PERP_RACE,
     family = binomial, data = clean_incident_df, na.action = na.exclude)
```

Coefficients:

	Estimate	Std. Error	z value	Pr(> z )	
(Intercept)	-1.06855	0.20611	-5.184	2.17e-07	***
BOROBROOKLYN	0.16083	0.05205	3.090	0.002003	**
BOROMANHATTAN	0.11455	0.06716	1.706	0.088079	.
BOROQUEENS	0.17563	0.06462	2.718	0.006571	**
BOROSTATEN ISLAND	0.22120	0.11703	1.890	0.058742	.
VIC_RACEAM_INDIAN/ALASKAN	-0.60707	1.05697	-0.574	0.565729	
VIC_RACEASIAN_PAC_ISLAND	-0.71881	0.19994	-3.595	0.000324	***
VIC_RACEBLACK	-0.62712	0.11250	-5.574	2.49e-08	***
VIC_RACEBLACK HISPANIC	-0.56914	0.12821	-4.439	9.04e-06	***
VIC_RACEUNKNOWN	-0.20656	0.37435	-0.552	0.581100	
VIC_RACEWHITE HISPANIC	-0.30462	0.11912	-2.557	0.010549	*
PERP_SEXM	-0.54537	0.12797	-4.262	2.03e-05	***
PERP_SEXU	-0.48918	0.23420	-2.089	0.036736	*
PERP_RACE(null)	-0.11578	0.27836	-0.416	0.677451	
PERP_RACEAM_INDIAN/ALASKAN	-8.69531	84.17442	-0.103	0.917724	
PERP_RACEASIAN_PAC_ISLAND	0.20867	0.28233	0.739	0.459852	
PERP_RACEBLACK	0.01052	0.17463	0.060	0.951985	
PERP_RACEBLACK HISPANIC	-0.32055	0.19623	-1.634	0.102363	
PERP_RACEUNKNOWN	-0.50673	0.25938	-1.954	0.050742	.
PERP_RACEWHITE HISPANIC	-0.25574	0.18346	-1.394	0.163326	

---

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 18197 on 28561 degrees of freedom  
Residual deviance: 17987 on 28542 degrees of freedom  
AIC: 18027

Number of Fisher Scoring iterations: 9

---

## Potential Bias

There are several sources of potential bias in this report



- Collection bias - Some of the categorical variable options, such as the descriptions of locations or locations categories, are limited and some actual locations may be ambiguous. The racial categorization options are also somewhat limited compared to the true diversity of racial and ethnic backgrounds. The racial categorization is likely to have been chosen by someone other than the subject, so the selected race may not be accurate or reflect how the subject (victim or perpetrator) would self-identify.
- Analysis bias - There are likely many interesting insights in this data regarding race (relationship between race of perpetrator and victim, racial distribution of incidents throughout the boroughs, etc), I deliberately avoided any in-depth analysis of those categories as I believed it would be more complicated and nuanced than my current ability and available time would allow me to do well.

```
sessionInfo()
```

```
R version 4.4.1 (2024-06-14)
Platform: aarch64-apple-darwin20
Running under: macOS Sonoma 14.6.1
```

```
Matrix products: default
```

```
BLAS: /Library/Frameworks/R.framework/Versions/4.4-arm64/Resources/lib/libRblas.0.dylib
```

```
LAPACK: /Library/Frameworks/R.framework/Versions/4.4-arm64/Resources/lib/libRlapack.dylib;
```

```
locale:
```

```
[1] C
```

```
time zone: America/Denver
```

```
tzcode source: internal
```

```
attached base packages:
```

```
[1] stats      graphics  grDevices  utils      datasets  methods   base
```

```
other attached packages:
```

```
[1] broom_1.0.6      forecast_8.23.0 lubridate_1.9.3 forcats_1.0.0
```

```
[5] stringr_1.5.1    dplyr_1.1.4      purrr_1.0.2     readr_2.1.5
```

```
[9] tidyr_1.3.1      tibble_3.2.1     ggplot2_3.5.1   tidyverse_2.0.0
```

```
loaded via a namespace (and not attached):
```

```
[1] utf8_1.2.4        generics_0.1.3    stringi_1.8.4    lattice_0.22-6
```

```
[5] hms_1.1.3         digest_0.6.36     magrittr_2.0.3    evaluate_0.24.0
```

```
[9] grid_4.4.1        timechange_0.3.0  pbdZMQ_0.3-11     fastmap_1.2.0
```

```
[13] jsonlite_1.8.8    backports_1.5.0   nnet_7.3-19       fansi_1.0.6
```

```
[17] scales_1.3.0      cli_3.6.3         rlang_1.1.4       crayon_1.5.3
```

[21]	munsell_0.5.1	base64enc_0.1-3	withr_3.0.1	repr_1.1.7
[25]	tools_4.4.1	parallel_4.4.1	tzdb_0.4.0	uuid_1.2-0
[29]	colorspace_2.1-1	curl_5.2.2	IRdisplay_1.1	vctrs_0.6.5
[33]	R6_2.5.1	zoo_1.8-12	lifecycle_1.0.4	tseries_0.10-58
[37]	urca_1.3-4	pkgconfig_2.0.3	pillar_1.9.0	gtable_0.3.5
[41]	quantmod_0.4.26	glue_1.7.0	Rcpp_1.0.13	lmtest_0.9-40
[45]	tidyselect_1.2.1	IRkernel_1.3.2	farver_2.1.2	nlme_3.1-164
[49]	htmltools_0.5.8.1	labeling_0.4.3	xts_0.14.0	timeDate_4041.110
[53]	fracdiff_1.5-3	compiler_4.4.1	quadprog_1.5-8	TTR_0.24.4