

NY Shooting Incident Report

CS

Assignment.

Import, tidy and analyze the NYPD Shooting Incident dataset obtained. Be sure your project is reproducible and contains some visualization and analysis. You may use the data to do any analysis that is of interest to you. You should include at least two visualizations and one model. Be sure to identify any bias possible in the data and in your analysis.

```
library(tidyverse)
```

```
-- Attaching core tidyverse packages ----- tidyverse 2.0.0 --
v dplyr      1.1.4      v readr      2.1.5
v forcats    1.0.0      v stringr    1.5.1
v ggplot2    3.5.1      v tibble     3.2.1
v lubridate  1.9.3      v tidyr      1.3.1
v purrr      1.0.2
-- Conflicts ----- tidyverse_conflicts() --
x dplyr::filter() masks stats::filter()
x dplyr::lag()     masks stats::lag()
i Use the conflicted package (<http://conflicted.r-lib.org/>) to force all conflicts to become
```

```
library(lubridate)
library(forecast)
```

```
Registered S3 method overwritten by 'quantmod':
  method      from
as.zoo.data.frame zoo
```

```
# import the source data and put it in a df
source_url <- paste0(
  "https://data.cityofnewyork.us/api/views/833y-fsy8/rows.csv?",
  "accessType=DOWNLOAD"
)

incident_df <- read.csv(source_url)
```

Explore

```
# see what columns we have and what data types
glimpse(incident_df)
```

```
Rows: 28,562
Columns: 21
$ INCIDENT_KEY      <int> 244608249, 247542571, 84967535, 202853370, 270~
$ OCCUR_DATE        <chr> "05/05/2022", "07/04/2022", "05/27/2012", "09/~
$ OCCUR_TIME        <chr> "00:10:00", "22:20:00", "19:35:00", "21:00:00"~
$ BORO              <chr> "MANHATTAN", "BRONX", "QUEENS", "BRONX", "BROO~
$ LOC_OF_OCCUR_DESC <chr> "INSIDE", "OUTSIDE", "", "", "", "", "", "", ""~
$ PRECINCT          <int> 14, 48, 103, 42, 83, 23, 113, 77, 48, 49, 73, ~
$ JURISDICTION_CODE <int> 0, 0, 0, 0, 0, 2, 0, 0, 0, 0, 0, 0, 0, 0, 0~
$ LOC_CLASSFCTN_DESC <chr> "COMMERCIAL", "STREET", "", "", "", "", "", ""~
$ LOCATION_DESC     <chr> "VIDEO STORE", "(null)", "", "", "", "MULTI DW~
$ STATISTICAL_MURDER_FLAG <chr> "true", "true", "false", "false", "false", "fa~
$ PERP_AGE_GROUP    <chr> "25-44", "(null)", "", "25-44", "25-44", "", "~
$ PERP_SEX          <chr> "M", "(null)", "", "M", "M", "", "", "", "", "~
$ PERP_RACE         <chr> "BLACK", "(null)", "", "UNKNOWN", "BLACK", "",~
$ VIC_AGE_GROUP     <chr> "25-44", "18-24", "18-24", "25-44", "25-44", "~
$ VIC_SEX           <chr> "M", "M", "M", "M", "M", "M", "M", "M", "M", "~
$ VIC_RACE          <chr> "BLACK", "BLACK", "BLACK", "BLACK", "BLACK", "~
$ X_COORD_CD        <dbl> 986050, 1016802, 1048632, 1014493, 1009149, 99~
$ Y_COORD_CD        <dbl> 214231.0, 250581.0, 198262.0, 242565.0, 190104~
$ Latitude          <dbl> 40.75469, 40.85440, 40.71063, 40.83242, 40.688~
$ Longitude         <dbl> -73.99350, -73.88233, -73.76777, -73.89071, -7~
$ Lon_Lat           <chr> "POINT (-73.9935 40.754692)", "POINT (-73.8823~
```

I already see lots of nulls, empty strings, missing values, etc. Let's take a closer look at some of the categorical columns to see if there are a limited number of consistently entered values or if they were entered as free text, which might be too difficult to clean.

```
# break out all the unique values and counts
desc_counts <- lapply(incident_df[,
  c(
    "LOC_CLASSFCTN_DESC",
    "LOCATION_DESC",
    "PERP_RACE",
    "VIC_RACE",
    "LOC_OF_OCCUR_DESC",
    "VIC_SEX",
    "PERP_SEX"
  )
], table)

print(desc_counts)
```

\$LOC_CLASSFCTN_DESC

	(null)	COMMERCIAL	DWELLING	HOUSING	OTHER
25596	2	208	243	460	59
PARKING LOT	PLAYGROUND	STREET	TRANSIT	VEHICLE	
15	41	1886	23	29	

\$LOCATION_DESC

	(null)	ATM
14977	1711	1
BANK	BAR/NIGHT CLUB	BEAUTY/NAIL SALON
3	668	119
CANDY STORE	CHAIN STORE	CHECK CASH
7	7	1
CLOTHING BOUTIQUE	COMMERCIAL BLDG	DEPT STORE
14	304	9
DOCTOR/DENTIST	DRUG STORE	DRY CLEANER/LAUNDRY
1	14	32
FACTORY/WAREHOUSE	FAST FOOD	GAS STATION
8	130	74
GROCERY/BODEGA	GYM/FITNESS FACILITY	HOSPITAL
750	4	77
HOTEL/MOTEL	JEWELRY STORE	LIQUOR STORE
35	14	42
LOAN COMPANY	MULTI DWELL - APT BUILD	MULTI DWELL - PUBLIC HOUS
1	2964	5007

NONE	PHOTO/COPY STORE	PVT HOUSE
175	1	983
RESTAURANT/DINER	SCHOOL	SHOE STORE
212	1	10
SMALL MERCHANT SOCIAL CLUB/POLICY LOCATI		STORAGE FACILITY
44	73	1
STORE UNCLASSIFIED	SUPERMARKET	TELECOMM. STORE
37	21	11
VARIETY STORE	VIDEO STORE	
11	8	

\$PERP_RACE

	(null)
9310	1141
AMERICAN INDIAN/ALASKAN NATIVE	ASIAN / PACIFIC ISLANDER
2	169
BLACK	BLACK HISPANIC
11903	1392
UNKNOWN	WHITE
1837	298
WHITE HISPANIC	
2510	

\$VIC_RACE

AMERICAN INDIAN/ALASKAN NATIVE	ASIAN / PACIFIC ISLANDER
11	440
BLACK	BLACK HISPANIC
20235	2795
UNKNOWN	WHITE
70	728
WHITE HISPANIC	
4283	

\$LOC_OF_OCCUR_DESC

	INSIDE	OUTSIDE
25596	460	2506

\$VIC_SEX

F	M	U
---	---	---

2760 25790 12

\$PERP_SEX

	(null)	F	M	U
9310	1141	444	16168	1499

Everything seems to be consistently entered (no misspellings or variations.) But there is a weird mix of “unknown”, “U”, and “null”. It will probably be best to recode empty values as “Unknown” for consistency. There is something weird in a few columns too.

```
unique(incident_df$PERP_RACE)
```

```
[1] "BLACK"                "(null)"
[3] ""                     "UNKNOWN"
[5] "WHITE HISPANIC"       "BLACK HISPANIC"
[7] "ASIAN / PACIFIC ISLANDER" "WHITE"
[9] "AMERICAN INDIAN/ALASKAN NATIVE"
```

```
table(incident_df$PERP_RACE)
```

	(null)
9310	1141
AMERICAN INDIAN/ALASKAN NATIVE	ASIAN / PACIFIC ISLANDER
2	169
BLACK	BLACK HISPANIC
11903	1392
UNKNOWN	WHITE
1837	298
WHITE HISPANIC	
2510	

Oh, that’s annoying - there is an empty string ' ' as one of the largest groups, I guess the best option will be to categorize that as "UNKNOWN" as well. While I’m at it I’m going to make the date time columns a little more usable by separating out the date and time and converting them to the right type.

Cleanup

```
clean_incident_df <- incident_df %>%
  mutate(
    # Combine the date and time into a proper DateTime object
    Date = as.POSIXct(
      paste(OCCUR_DATE, OCCUR_TIME),
      format = "%m/%d/%Y %H:%M:%S"
    )
  ) %>%
  rename(
    In_Out = LOC_OF_OCCUR_DESC,
    Location_Category = LOC_CLASSFCTN_DESC,
    Location_details = LOCATION_DESC
  ) %>%
  select(
    Date, BORO, Location_Category, Location_details,
    In_Out, OCCUR_DATE, OCCUR_TIME,
    -JURISDICTION_CODE, -X_COORD_CD, -Y_COORD_CD,
    -Latitude, -Longitude, -Lon_Lat, -PRECINCT,
    everything()
  ) %>%
  mutate(
    # Recode specific values in PERP_RACE and VIC_RACE
    PERP_RACE = recode(PERP_RACE,
      "ASIAN / PACIFIC ISLANDER" = "ASIAN_PAC_ISLAND",
      "AMERICAN INDIAN/ALASKAN NATIVE" = "AM_INDIAN/ALASKAN"
    ),
    VIC_RACE = recode(VIC_RACE,
      "ASIAN / PACIFIC ISLANDER" = "ASIAN_PAC_ISLAND",
      "AMERICAN INDIAN/ALASKAN NATIVE" = "AM_INDIAN/ALASKAN"
    ),
    # Recode empty or null values to "UNKNOWN" for PERP_RACE
    PERP_RACE = ifelse(PERP_RACE == "", "UNKNOWN", PERP_RACE),

    # Recode unknown, empty, or NA values to "U" for both victim and perpetrator sex
    PERP_SEX = ifelse(PERP_SEX %in% c("Unknown", "", (null), NA), "U", PERP_SEX),
    VIC_SEX = ifelse(VIC_SEX %in% c("Unknown", "", NA), "U", VIC_SEX)
  )

# Check that the recoding worked as expected
unique(clean_incident_df$PERP_SEX) # Should show only "M", "F", and "U"
```

```
[1] "M" "U" "F"
```

```
unique(clean_incident_df$VIC_SEX) # Should show only "M", "F", and "U"
```

```
[1] "M" "F" "U"
```

```
# check that I have the columns and order that I wanted  
glimpse(clean_incident_df)
```

```
Rows: 28,562
```

```
Columns: 22
```

```
$ Date          <dtm> 2022-05-05 00:10:00, 2022-07-04 22:20:00, 201~  
$ BORO          <chr> "MANHATTAN", "BRONX", "QUEENS", "BRONX", "BROO~  
$ Location_Category <chr> "COMMERCIAL", "STREET", "", "", "", "", "", ""~  
$ Location_details <chr> "VIDEO STORE", "(null)", "", "", "", "MULTI DW~  
$ In_Out        <chr> "INSIDE", "OUTSIDE", "", "", "", "", "", ""~  
$ OCCUR_DATE    <chr> "05/05/2022", "07/04/2022", "05/27/2012", "09/~  
$ OCCUR_TIME    <chr> "00:10:00", "22:20:00", "19:35:00", "21:00:00"~  
$ INCIDENT_KEY  <int> 244608249, 247542571, 84967535, 202853370, 270~  
$ PRECINCT      <int> 14, 48, 103, 42, 83, 23, 113, 77, 48, 49, 73, ~  
$ JURISDICTION_CODE <int> 0, 0, 0, 0, 0, 2, 0, 0, 0, 0, 0, 0, 0, 0, 0~  
$ STATISTICAL_MURDER_FLAG <chr> "true", "true", "false", "false", "false", "fa~  
$ PERP_AGE_GROUP <chr> "25-44", "(null)", "", "25-44", "25-44", "", "~  
$ PERP_SEX      <chr> "M", "U", "U", "M", "M", "U", "U", "U", "U", "~  
$ PERP_RACE     <chr> "BLACK", "(null)", "UNKNOWN", "UNKNOWN", "BLAC~  
$ VIC_AGE_GROUP <chr> "25-44", "18-24", "18-24", "25-44", "25-44", "~  
$ VIC_SEX       <chr> "M", "M", "M", "M", "M", "M", "M", "M", "M", "~  
$ VIC_RACE      <chr> "BLACK", "BLACK", "BLACK", "BLACK", "BLACK", "~  
$ X_COORD_CD    <dbl> 986050, 1016802, 1048632, 1014493, 1009149, 99~  
$ Y_COORD_CD    <dbl> 214231.0, 250581.0, 198262.0, 242565.0, 190104~  
$ Latitude      <dbl> 40.75469, 40.85440, 40.71063, 40.83242, 40.688~  
$ Longitude     <dbl> -73.99350, -73.88233, -73.76777, -73.89071, -7~  
$ Lon_Lat       <chr> "POINT (-73.9935 40.754692)", "POINT (-73.8823~
```

```
# check that we fixed the empty string values  
unique(clean_incident_df$PERP_RACE)
```

```
[1] "BLACK"          "(null)"          "UNKNOWN"
[4] "WHITE HISPANIC" "BLACK HISPANIC"  "ASIAN_PAC_ISLAND"
[7] "WHITE"          "AM_INDIAN/ALASKAN"
```

```
# printing a df is a little uglier in some ways but prevents text
# overlap when there are lots of columns or long column names
print(tail(clean_incident_df))
```

	Date	BORO	Location_Category	Location_details		
28557	2023-07-02 21:40:00	BRONX	STREET	(null)		
28558	2023-03-19 23:48:00	BRONX	COMMERCIAL	GROCERY/BODEGA		
28559	2023-08-16 02:46:00	BRONX	STREET	(null)		
28560	2023-06-27 12:27:00	BRONX	DWELLING	MULTI DWELL - APT BUILD		
28561	2023-07-08 11:27:00	QUEENS	STREET	BEAUTY/NAIL SALON		
28562	2023-07-24 23:38:00	MANHATTAN	HOUSING	MULTI DWELL - PUBLIC HOUS		
	In_Out	OCCUR_DATE	OCCUR_TIME	INCIDENT_KEY	PRECINCT	JURISDICTION_CODE
28557	OUTSIDE	07/02/2023	21:40:00	270719378	46	0
28558	INSIDE	03/19/2023	23:48:00	265354835	47	0
28559	OUTSIDE	08/16/2023	02:46:00	272968931	41	0
28560	INSIDE	06/27/2023	12:27:00	270489846	41	0
28561	OUTSIDE	07/08/2023	11:27:00	271021661	102	0
28562	OUTSIDE	07/24/2023	23:38:00	271818283	28	2
	STATISTICAL_MURDER_FLAG	PERP_AGE_GROUP	PERP_SEX	PERP_RACE		
28557	false	(null)	U	(null)		
28558	true	18-24	M	BLACK		
28559	false	25-44	F	BLACK		
28560	true	25-44	M	BLACK		
28561	false	25-44	M	WHITE HISPANIC		
28562	false	(null)	U	(null)		
	VIC_AGE_GROUP	VIC_SEX	VIC_RACE	X_COORD_CD	Y_COORD_CD	Latitude
28557	18-24	M	BLACK HISPANIC	1009601	247515	40.84601
28558	18-24	M	BLACK	1025687	268586	40.90378
28559	45-64	M	BLACK	1014639	240066	40.82555
28560	25-44	M	BLACK	1012221	238552	40.82140
28561	65+	M	ASIAN_PAC_ISLAND	1028856	192785	40.69572
28562	25-44	M	BLACK	997853	230889	40.80040
	Longitude	Lon_Lat				
28557	-73.90837	POINT (-73.908369 40.846012)				
28558	-73.85010	POINT (-73.850098 40.903785)				
28559	-73.89020	POINT (-73.890195 40.825549)				
28560	-73.89894	POINT (-73.898938 40.821404)				
28561	-73.83914	POINT (-73.839138 40.695717)				

28562 -73.95086 POINT (-73.950864 40.800405)

I'm going to make a few different dataframes with different groups for eventual analysis and plotting. Things I'm going to start with

- Daily incidents over time to look for general trends
- Incidents by month and year
- Incidents by borough
- Incidents by month (not over time, so total incidents that occurred in each month summed over all years)
- Time and year data broken down by borough
- Victim and perpetrator by sex

```
# for plotting incidents over time
time_series_df <- clean_incident_df %>%
  mutate(simple_date = as.Date(OCCUR_DATE, format = "%m/%d/%Y")) %>%
  group_by(simple_date) %>%
# Add a new column that represents only the month and year
# This step may be unnecessary since I have a good date column
# but it's easier for me to understand
  summarise(total_by_day = n()) %>%
  mutate(month_year = floor_date(simple_date, "month"))

# for plotting overtime by month and year
df_aggregated <- time_series_df %>%
  mutate(year = format(simple_date, "%Y"),
         month = format(simple_date, "%m")) %>%
  group_by(year, month) %>%
  summarise(total_by_month = sum(total_by_day)) %>%
  mutate(
    year = as.numeric(year), # Convert year to numeric
    month = as.numeric(month) # Convert month to numeric
  ) %>%
  ungroup()
```

`summarise()` has grouped output by 'year'. You can override using the
`.groups` argument.

```
# borough totals
total_by_borough <- clean_incident_df %>%
  group_by(BORO) %>%
  summarize(total_incidents = n())
```

```
# monthly borough totals
monthly_totals_by_borough <- clean_incident_df %>%
  mutate(month = floor_date(Date, "month")) %>%
  mutate(month = as.Date(month)) %>%
  group_by(BORO, month) %>%
  summarize(monthly_incidents = n()) %>%
  ungroup()
```

`summarise()` has grouped output by 'BORO'. You can override using the
`.groups` argument.

```
# victim and perp by sex
totals_by_sex <- clean_incident_df %>%
  group_by(PERP_SEX, VIC_SEX) %>%
  summarise(
    Total_Victims = n(),      # Count number of victims in each group
    Total_Perps = n()        # Count number of perpetrators in each group
  )
```

`summarise()` has grouped output by 'PERP_SEX'. You can override using the
`.groups` argument.

```
tail(time_series_df)
```

```
# A tibble: 6 x 3
  simple_date total_by_day month_year
  <date>         <int> <date>
1 2023-12-22         8 2023-12-01
2 2023-12-23         4 2023-12-01
3 2023-12-24         5 2023-12-01
4 2023-12-26         6 2023-12-01
5 2023-12-27         1 2023-12-01
6 2023-12-29         3 2023-12-01
```

```
tail(df_aggregated)
```

```
# A tibble: 6 x 3
  year month total_by_month
  <dbl> <dbl>         <int>
1 2023 12         26
2 2023 12         19
3 2023 12         14
4 2023 12         10
5 2023 12          7
6 2023 12          3
```

1	2023	7	152
2	2023	8	108
3	2023	9	105
4	2023	10	99
5	2023	11	71
6	2023	12	83

```
tail(total_by_borough)
```

```
# A tibble: 5 x 2
  BORO      total_incidents
  <chr>          <int>
1 BRONX             8376
2 BROOKLYN          11346
3 MANHATTAN          3762
4 QUEENS             4271
5 STATEN ISLAND       807
```

```
tail(monthly_totals_by_borough)
```

```
# A tibble: 6 x 3
  BORO      month      monthly_incidents
  <chr>    <date>          <int>
1 STATEN ISLAND 2023-05-01         3
2 STATEN ISLAND 2023-06-01         8
3 STATEN ISLAND 2023-07-01         6
4 STATEN ISLAND 2023-08-01         3
5 STATEN ISLAND 2023-10-01         3
6 STATEN ISLAND 2023-11-01         2
```

```
tail(totals_by_sex,9)
```

```
# A tibble: 9 x 4
# Groups:   PERP_SEX [3]
  PERP_SEX VIC_SEX Total_Victims Total_Perps
  <chr>    <chr>          <int>    <int>
1 F      F             77         77
2 F      M            366        366
3 F      U             1          1
4 M      F           1755       1755
```

5	M	M	14406	14406
6	M	U	7	7
7	U	F	928	928
8	U	M	11018	11018
9	U	U	4	4

That's looks pretty good. We'll progress through some basic plots to see what patterns or trends we see which can guide the next visualizations.

Visualization

```
# Set global plot size options for the notebook
options(repr.plot.width = 10, repr.plot.height = 7)
```

```
theme_set(theme_minimal())

# plot daily incidents
ggplot(time_series_df, aes(x = simple_date, y = total_by_day)) +
  geom_line(color = "skyblue") +
  labs(
    title = "Fig.1 - Daily Incidents",
    x = "Date",
    y = "Incident Count"
  )
```

Fig.1 – Daily Incidents

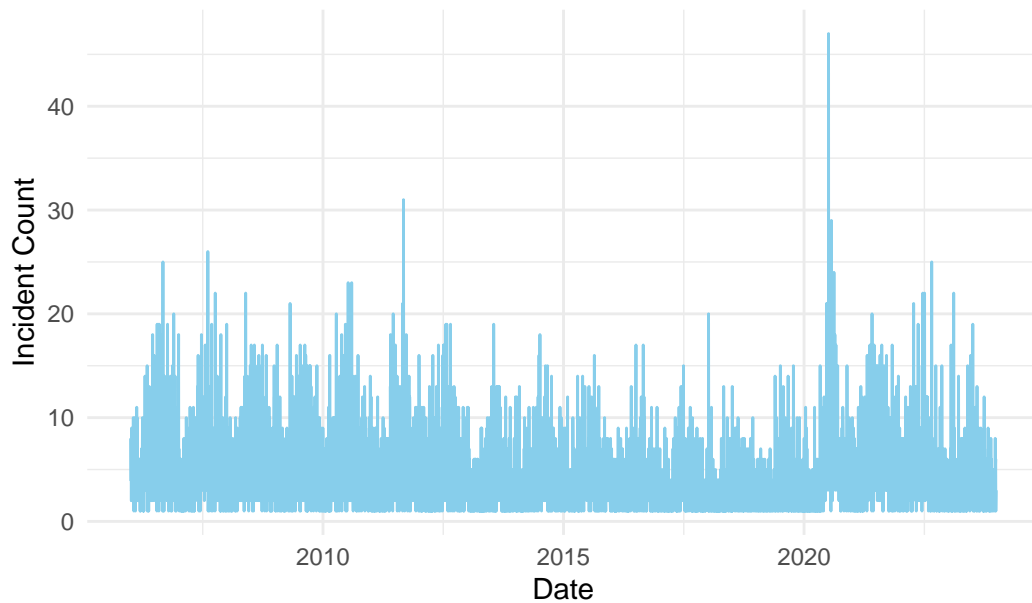
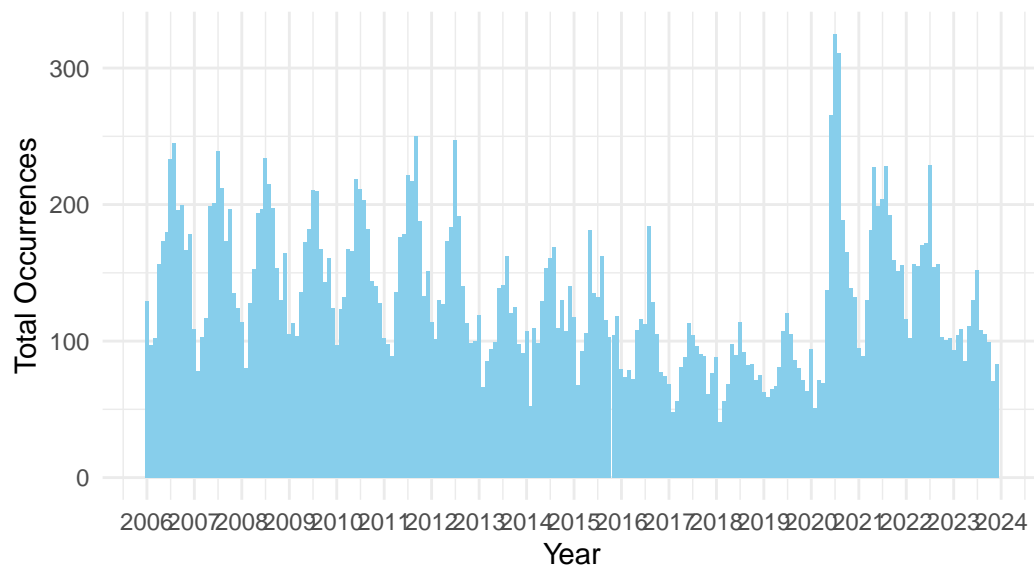


Fig 1 - Daily incidents - hard to see many trends since the data is so noisy, but it does look like it's generally periodic and there is a big spike around the first covid summer.

```
# plot monthly incidents over time
ggplot(time_series_df, aes(x = month_year, y = total_by_day)) +
  geom_bar(stat = "identity", fill = "skyblue") +
  labs(
    title = "Fig. 2 - Total Occurrences by Month",
    x = "Year",
    y = "Total Occurrences",
    caption = "Source: NYPD Shooting Incident Data, 2020") +
  scale_x_date(date_labels = "%Y", date_breaks = "1 year")
```

Fig. 2 – Total Occurrences by Month



Source: NYPD Shooting Incident Data, 2020

Fig 2 - Total occurrences by month over time - clearer seasonality and a little easier to see the pre- and post-covid trends

```
# plot yearly incidents
ggplot(time_series_df, aes(
  x = year(simple_date),
  y = total_by_day)) +
  geom_bar(stat = "identity", fill = "skyblue") +
  labs(
    title = "Fig. 3 - Total Occurrences by Year",
    x = "Year",
    y = "Total Occurrences")
```

Fig. 3 – Total Occurrences by Year

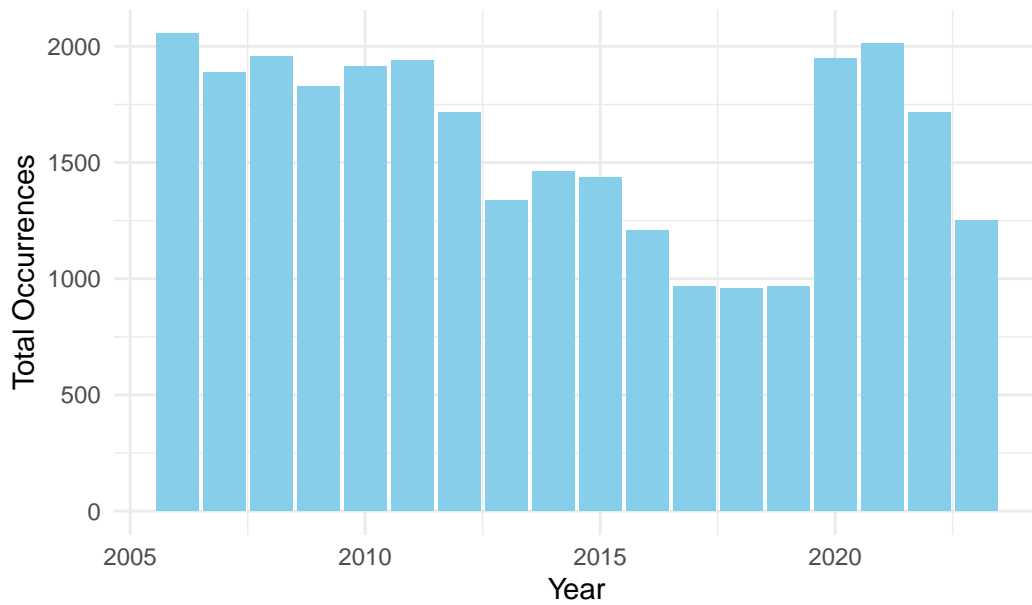


Fig 3 - Total occurrences by year - now we can see trends. Decreasing incidents starting in the early 20-teens and flattening out, before a big covid spike and almost back down to pre-covid levels.

```
# Plot occurrences by month (across all years)
ggplot(time_series_df, aes(
  x = month(simple_date, label = TRUE),
  y = total_by_day
)) +
  geom_bar(
    stat = "identity",
    fill = "skyblue"
  ) +
  labs(
    title = "Fig. 4 - Total Occurrences by Month",
    x = "Month",
    y = "Total Occurrences"
  )
```

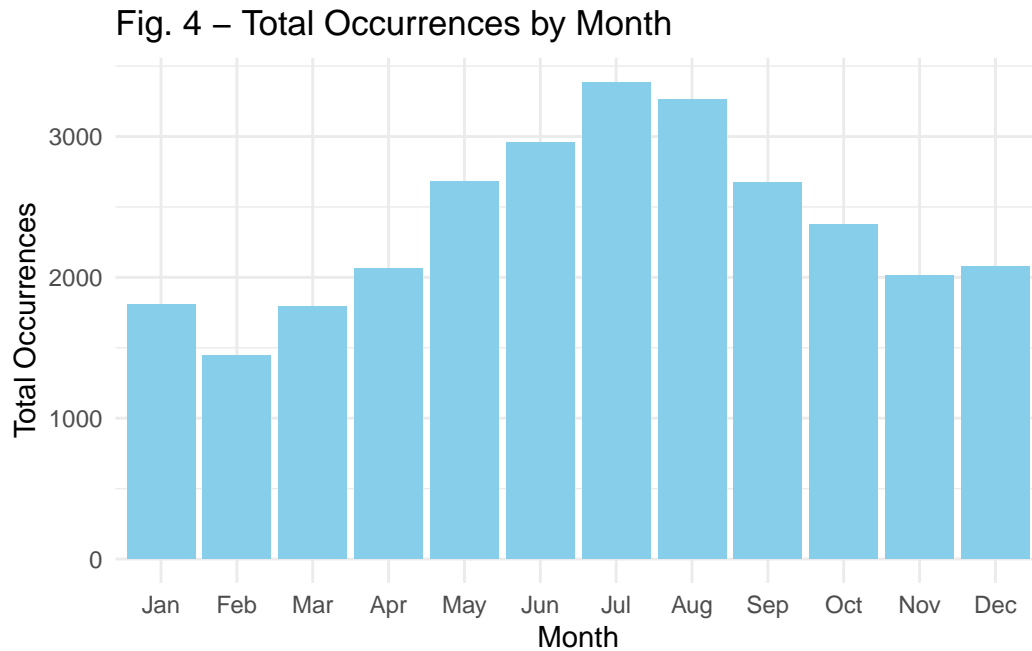


Fig 4 - Total occurrences by month - pretty strong visual trend towards higher incidents in the hottest months, which is a well studied phenomenon.

```
# Plot total by borough
ggplot(total_by_borough, aes(
  x = BORO,
  y = total_incidents
)) +
  geom_bar(
    stat = "identity",
    fill = "skyblue"
  ) +
  labs(
    title = "Fig. 5 - Total by Borough",
    x = "Borough",
    y = "Total"
  )
```

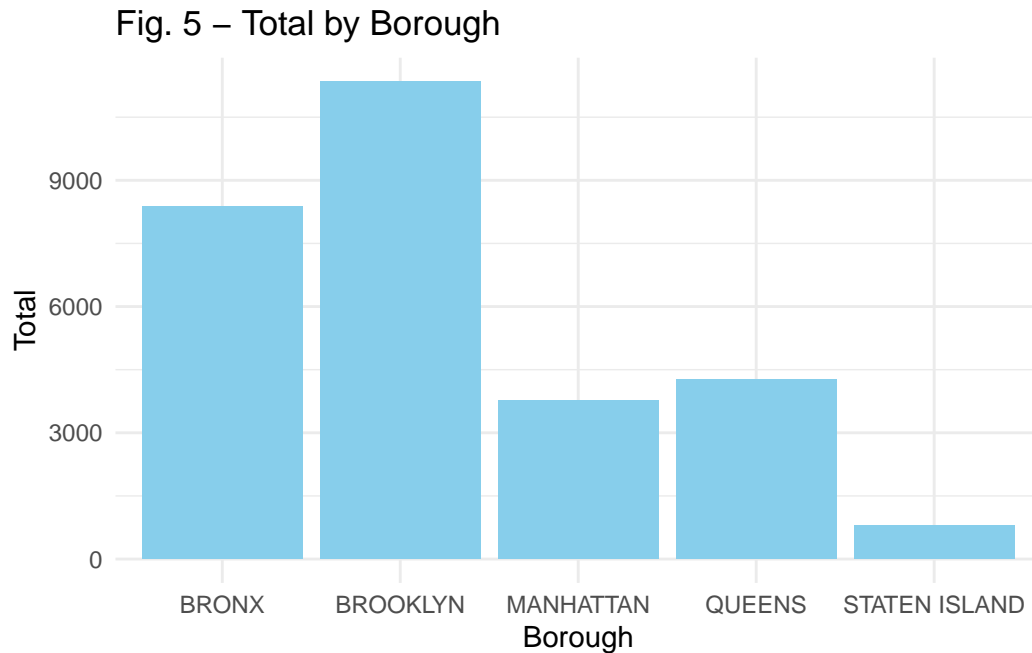
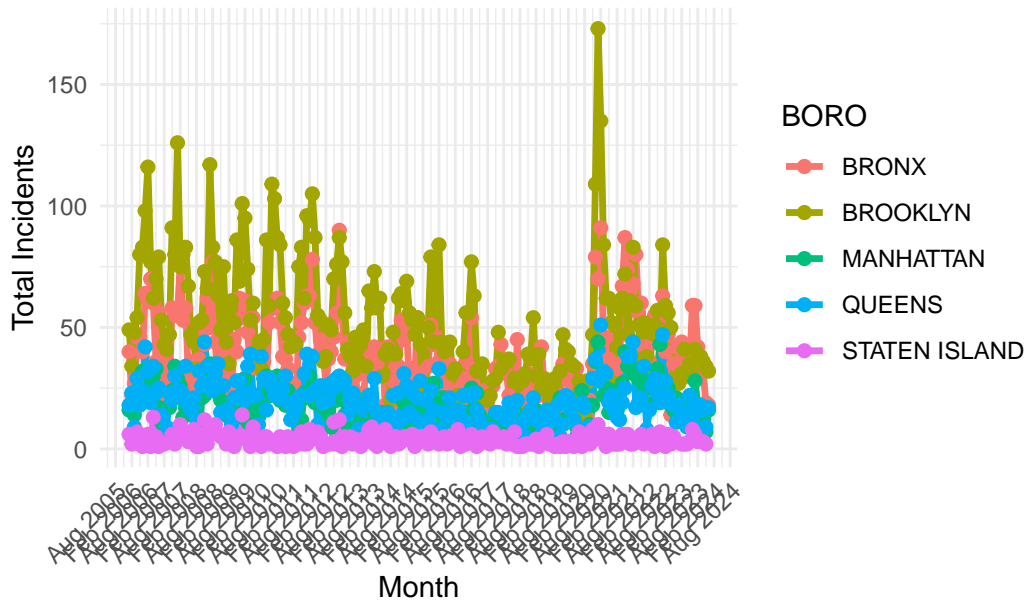



Fig 5 - Total by borough - generally interesting, but would be more useful with the context of per capita and per area data for the boroughs.

```
# Borough totals monthly
ggplot(monthly_totals_by_borough, aes(
  x = month,
  y = monthly_incidents,
  color = BORO)
) +
# I wanted to try a line instead of a bar
geom_line(linewidth = 1.2) +
geom_point(size = 2) +
labs(
  title = "Fig. 6 - Monthly Incidents by Borough",
  x = "Month",
  y = "Total Incidents"
) +
scale_x_date(date_labels = "%b %Y", date_breaks = "6 month") +
theme(axis.text.x = element_text(angle = 45, hjust = 1))
```

Fig. 6 – Monthly Incidents by Borough



```
# Plot incidents by year for each borough with side-by-side bars
ggplot(monthly_totals_by_borough, aes(
  x = year(month),
  y = monthly_incidents,
  fill = BORO)
) +
# Use dodge for side-by-side bars
geom_bar(stat = "identity", position = "dodge") +
labs(
  title = "Fig. 7 - Yearly Incidents by Borough",
  x = "Year",
  y = "Total Incidents")
```

Fig. 7 – Yearly Incidents by Borough

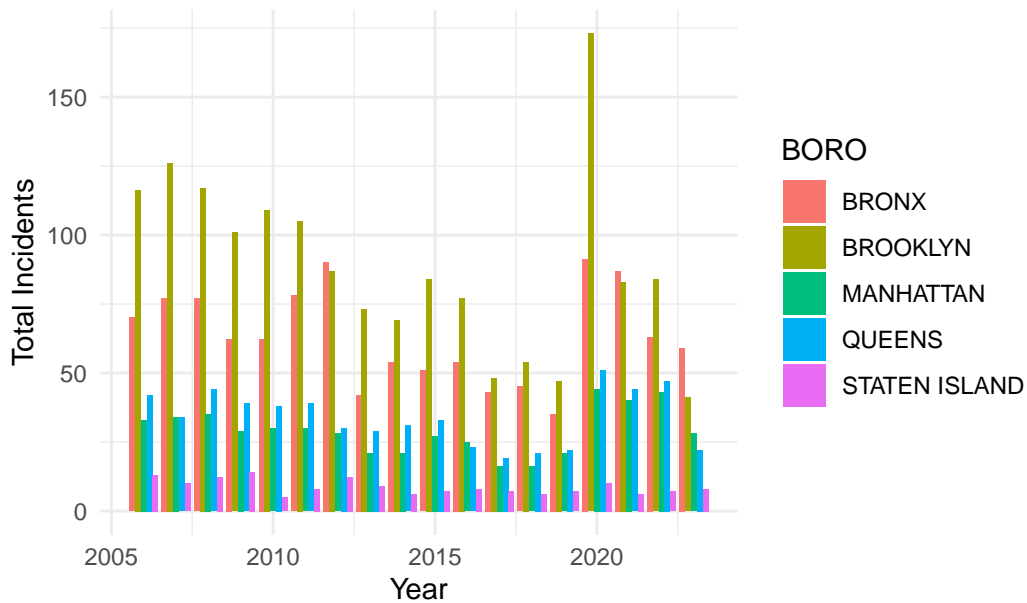


Fig 6/7 - Boroughs over time - interesting to see that not all boroughs follow the same trends over time, and that the first covid spike was driven heavily by increased in Brooklyn.

```
# plot for total victims of each sex by perpetrator sex
ggplot(totals_by_sex, aes(
  x = PERP_SEX,
  y = Total_Victims,
  fill = VIC_SEX)
) +
geom_bar(stat = "identity", position = "dodge") +
labs(
  title = "Fig. 8 - Victim Sex by Perpetrator Sex",
  x = "Perpetrator Sex",
  y = "Total Victims",
  fill = "Victim Sex"
) +
theme_minimal()
```

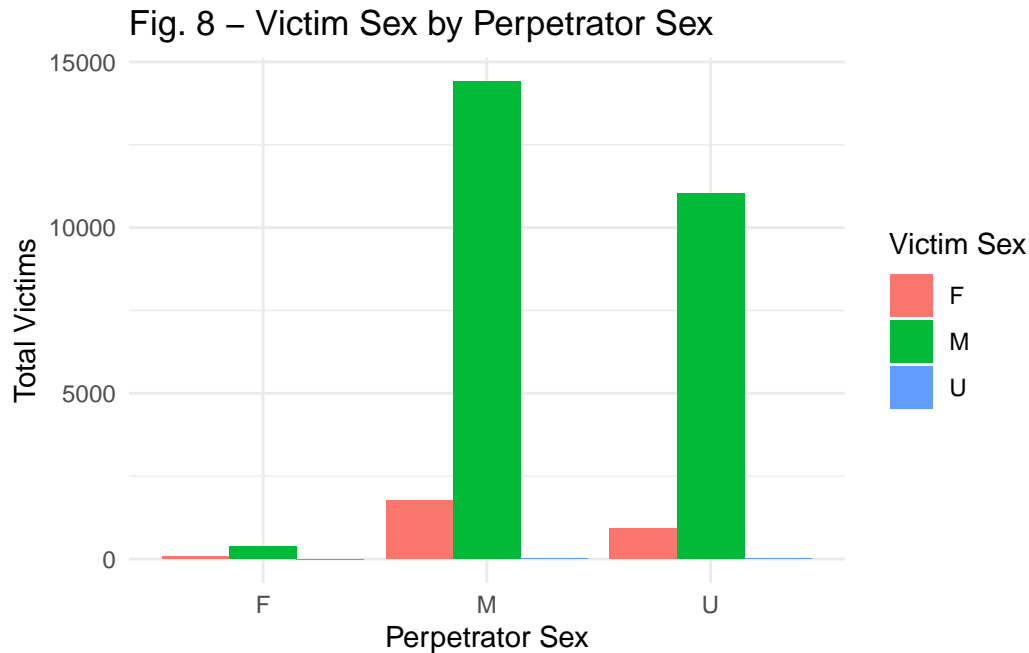


Fig 8 - Totals of victim sex group by perpetrator sex - there are fewer total female victims than I would have expected. I'm going to use this visualization as inspiration for my first simple model below.

We could keep going with similar visuals (breakdown by race, gender, age group, etc. or relationships between group like victim age relative to perpetrator age) but I'll stop there. I'm going to focus on victim sex for the analysis and modeling component. I want to see how predictive of a victim being female some of the other attributes (perpetrator race and sex, victim race). I'll start with a simple model of victim sex as predicted by perpetrator sex. To start will have to exclude the "unknowns" from victim sex and clean up some of the other factors.

Modeling

Sex differences among victims and perpetrators

For my first simple model I want to do a basic breakdown or how a perpetrator of a given sex affects the odds of the victim being a give sex.

```
# Convert all categorical variables to factors
clean_incident_df$VIC_SEX <- factor(clean_incident_df$VIC_SEX)
clean_incident_df$BORO <- factor(clean_incident_df$BORO)
clean_incident_df$VIC_RACE <- factor(clean_incident_df$VIC_RACE)
clean_incident_df$PERP_RACE <- factor(clean_incident_df$PERP_RACE)
clean_incident_df$PERP_SEX <- factor(clean_incident_df$PERP_SEX)

# Apply droplevels to all factor columns (to remove unused levels)
clean_incident_df <- clean_incident_df %>%
  mutate_if(is.factor, droplevels)

# Set "M" as the victim reference so that we model the odds of being "Female"
# Set the reference race for victims and perps as "white"
clean_incident_df$VIC_SEX <- relevel(clean_incident_df$VIC_SEX, ref = "M")
clean_incident_df$VIC_RACE <- relevel(clean_incident_df$VIC_RACE, ref = "WHITE")
clean_incident_df$PERP_RACE <- relevel(clean_incident_df$PERP_RACE, ref = "WHITE")

# Check the levels of the factor to confirm they are correct
levels(clean_incident_df$VIC_SEX)
```

```
[1] "M" "F" "U"
```

```
levels(clean_incident_df$PERP_SEX)
```

```
[1] "F" "M" "U"
```

```
levels(clean_incident_df$PERP_RACE)
```

```
[1] "WHITE"          "(null)"          "AM_INDIAN/ALASKAN"
[4] "ASIAN_PAC_ISLAND" "BLACK"           "BLACK HISPANIC"
[7] "UNKNOWN"         "WHITE HISPANIC"
```

```
# create a simple generalize linear model to predict odds of female victim
# based on perp sex
simple_model_vic_sex <- glm(VIC_SEX ~ PERP_SEX,
  family=binomial,
  data=clean_incident_df,
  na.action = na.exclude)

summary(simple_model_vic_sex)
```

```
Call:
glm(formula = VIC_SEX ~ PERP_SEX, family = binomial, data = clean_incident_df,
     na.action = na.exclude)
```

Coefficients:

	Estimate	Std. Error	z value	Pr(> z)
(Intercept)	-1.5459	0.1247	-12.396	< 2e-16 ***
PERP_SEXM	-0.5553	0.1272	-4.364	1.28e-05 ***
PERP_SEXU	-0.9240	0.1293	-7.147	8.88e-13 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 18197 on 28561 degrees of freedom
Residual deviance: 18093 on 28559 degrees of freedom
AIC: 18099

Number of Fisher Scoring iterations: 5

It looks like each perpetrator sex option has a statistically significant effect on the change in log-odds of the victim being female, so I'll pull out each of them and convert them to odds.

```
print("Log-odds are")
```

```
[1] "Log-odds are"
```

```
print(coef(simple_model_vic_sex)[c("PERP_SEXM", "PERP_SEXU")])
```

```
PERP_SEXM PERP_SEXU
-0.5552708 -0.9240283
```

```
#convert log-odds to odds and print
```

```
print("The odds relative to the female victim/female perpetrator baseline are")
```

```
[1] "The odds relative to the female victim/female perpetrator baseline are"
```

```
print((exp(coef(simple_model_vic_sex)[c("PERP_SEXM", "PERP_SEXU"]))))
```

```
PERP_SEXM PERP_SEXU  
0.5739169 0.3969169
```

```
# Print odds for when the perpetrator is female  
paste(  
  "The odds of a victim being female when the perpetrator is female are",  
  round(exp(coef(simple_model_vic_sex)["(Intercept)"]), 2)  
)
```

```
[1] "The odds of a victim being female when the perpetrator is female are 0.21"
```

```
# Print odds for when the perpetrator is male  
paste0(  
  "The odds of a victim being female when the perpetrator is male are ",  
  round(exp(coef(simple_model_vic_sex)["PERP_SEXM"]), 2) * 100,  
  "% of baseline"  
)
```

```
[1] "The odds of a victim being female when the perpetrator is male are 57% of baseline"
```

```
# Print odds for when the perpetrator is unknown  
paste0(  
  "The odds of a victim being female when the perpetrator is unknown are ",  
  round(exp(coef(simple_model_vic_sex)["PERP_SEXU"]), 2) * 100,  
  "% of baseline"  
)
```

```
[1] "The odds of a victim being female when the perpetrator is unknown are 40% of baseline"
```

Simple Summer prediction from Winter data

This isn't exactly a model, but I wanted to see how consistently the ratios of incidents in the Summer relative to Winter have been year over year.

```
# make a better df for monthly modeling
monthly_totals <- time_series_df %>%
  group_by(month_year) %>%
  summarise(total_by_month = sum(total_by_day))

# Check the aggregated data
tail(monthly_totals)
```

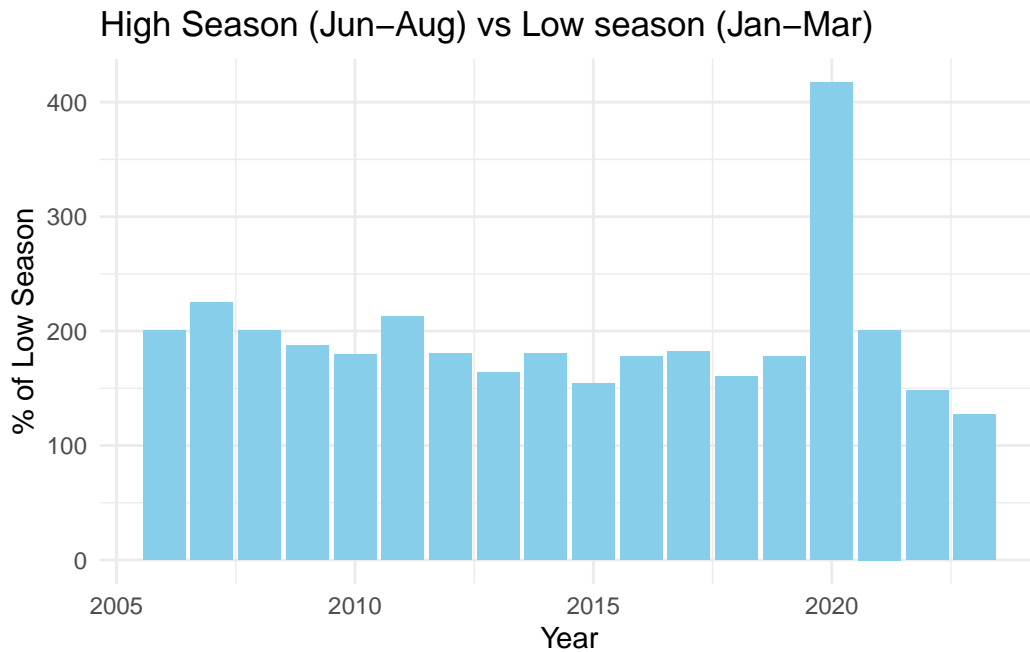
```
# A tibble: 6 x 2
  month_year total_by_month
  <date>         <int>
1 2023-07-01         152
2 2023-08-01         108
3 2023-09-01         105
4 2023-10-01          99
5 2023-11-01          71
6 2023-12-01          83
```

```
# make a simple ratio model based on high and low seasons
# Aggregate by year and calculate sum for January-March and
seasonal_totals <- df_aggregated %>%
  mutate(season = case_when(
    month %in% c(1, 2, 3) ~ "low_season",
    month %in% c(6, 7, 8) ~ "high_season",
    TRUE ~ NA_character_
  )) %>%
  filter(!is.na(season)) %>% # Filter to keep only rows with season
  group_by(year, season) %>%
  summarise(total = sum(total_by_month, na.rm = TRUE), .groups = "drop") %>%
  pivot_wider(names_from = season, values_from = total) %>%
  mutate(ratio = (high_season / low_season) * 100)

# plot the simple ratio results
ggplot(seasonal_totals, aes(
  x = year,
  y = ratio
)) +
  geom_bar(
    stat = "identity",
    fill = "skyblue"
  ) +
```



```
labs(
  title = "High Season (Jun-Aug) vs Low season (Jan-Mar)",
  x = "Year",
  y = "% of Low Season"
)
```



I think this is an interesting view of how consistently the ratio of incidents in the worst 3 months each year are to the lowest 3 months. This is an example of a quick and easy naive estimation that could be used for things like predicting staffing or resource needs for later in the year.

More advance forecasting model

I want to take it a step further and see what could be done with some simple forecasting models. Will use seasonal ARIMA (SARIMA) which considers seasonality as well as recent trends to create a future forecast. It's import to not that this model uses only the previous data, it doesn't consider other variables. For instance, if there are always higher shooting rates when the year before was hotter or the stock market was lower, it won't include that information.

```
# make time series friendly
start_year <- as.numeric(format(min(monthly_totals$month_year), "%Y"))
start_month <- as.numeric(format(min(monthly_totals$month_year), "%m"))

# Convert to a time series object (monthly frequency)
monthly_incident_ts <- ts(monthly_totals$total_by_month, start = c(start_year, start_month),

# Check the time series
print(monthly_incident_ts)
```

	Jan	Feb	Mar	Apr	May	Jun	Jul	Aug	Sep	Oct	Nov	Dec
2006	129	97	102	156	173	180	233	245	196	199	167	178
2007	109	78	103	117	199	201	239	212	173	197	135	124
2008	114	80	128	153	194	197	234	215	197	153	130	164
2009	105	113	104	136	172	182	211	210	167	143	161	124
2010	97	123	132	167	166	219	211	203	182	144	140	128
2011	102	98	89	136	176	178	221	217	250	188	133	151
2012	114	101	130	127	173	183	247	191	140	113	98	100
2013	119	66	85	94	99	139	141	162	120	125	98	91
2014	107	52	109	98	129	153	161	169	109	130	107	140
2015	117	68	93	106	181	135	132	162	115	103	104	118
2016	79	74	79	72	108	116	112	184	128	105	77	74
2017	68	48	56	81	88	113	104	96	90	89	61	76
2018	88	41	56	68	98	90	114	92	82	83	71	75
2019	63	59	65	67	81	107	120	105	86	80	71	63
2020	94	51	71	69	137	265	325	311	189	165	139	132
2021	95	89	130	181	227	199	204	228	192	159	151	156
2022	116	102	156	155	170	172	229	154	156	103	101	102
2023	93	104	109	85	111	130	152	108	105	99	71	83

```
sarima_forecast <- auto.arima(monthly_incident_ts, seasonal = TRUE)

# Summary of the SARIMA model
summary(sarima_forecast)
```

Series: monthly_incident_ts
 ARIMA(1,0,0)(1,1,0)[12] with drift

Coefficients:

ar1	sar1	drift
0.6957	-0.3703	-0.3007

s.e. 0.0507 0.0662 0.4204

$\sigma^2 = 906.9$: log likelihood = -983.79

AIC=1975.58 AICc=1975.78 BIC=1988.85

Training set error measures:

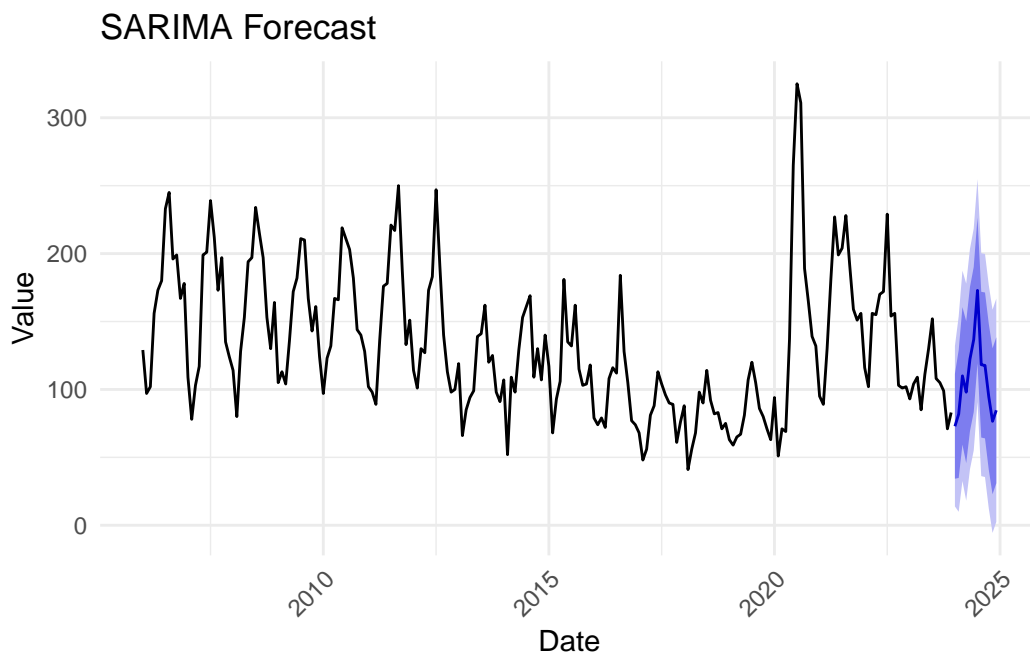
	ME	RMSE	MAE	MPE	MAPE	MASE
Training set	-0.01241811	29.05056	21.49203	-2.150574	17.37663	0.7560569

ACF1

Training set -0.07634969

```
# Forecasting the next 12 months
forecast_values <- forecast::forecast(sarima_forecast, h = 12)

# Plot the forecast
autoplot(forecast_values) +
  labs(x = "Date", y = "Value", title = "SARIMA Forecast") +
  theme(axis.text.x = element_text(angle = 45, hjust = 1))
```



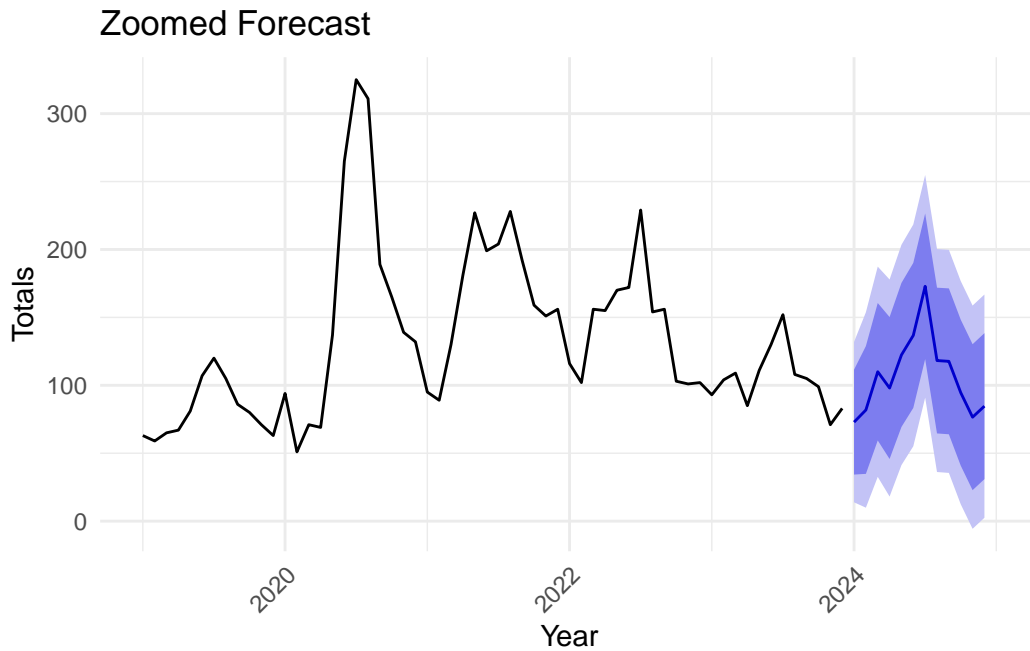
```
# create a zoomed in version
autoplot(forecast_values) +
  labs(title = "Zoomed Forecast", x = "Year", y = "Totals") +
```

```
xlim(c(2019, 2025)) + # Adjust x-axis to show the previous 12 months and forecast
theme(axis.text.x = element_text(angle = 45, hjust = 1))
```

Scale for x is already present.

Adding another scale for x, which will replace the existing scale.

Warning: Removed 156 rows containing missing values or values outside the scale range (`geom_line()`).



```
# Print the confidence intervals for the next 5 months
next_12_conf_intervals <- data.frame(
  month2024 = c(1:12),
  forecast = forecast_values$mean[1:12],
  lower_80 = forecast_values$lower[1:12, 1], # 80% lower bound
  upper_80 = forecast_values$upper[1:12, 1], # 80% upper bound
  lower_95 = forecast_values$lower[1:12, 2],
  upper_95 = forecast_values$upper[1:12, 2]
) %>%
  mutate(across(everything(), \(x) round(x, 0)))
next_12_conf_intervals
```

	month2024	forecast	lower_80	upper_80	lower_95	upper_95
1	1	73	34	111	14	132
2	2	82	35	129	10	154
3	3	110	59	161	33	187
4	4	98	46	150	18	178
5	5	122	69	175	41	203
6	6	137	83	190	55	218
7	7	173	119	226	91	255
8	8	118	65	172	36	200
9	9	118	64	171	36	200
10	10	95	41	148	12	177
11	11	77	23	130	-6	159
12	12	85	31	138	2	167

That's pretty neat. Relatively wide confidence intervals but the forecasted trend looks reasonable and could be useful for relatively low effort.

That's the end of the analysis I'm comfortable with. Below I wanted to see what it would look like to do similar modeling with multiple predictor variables (perp sex, race, borough.) It looks like it worked, but it gets out of hand to interpret it pretty quickly so I just stopped and left it here as an interesting example of what else could be done.

```
model_vic_sex <- glm(VIC_SEX ~ BORO + VIC_RACE +
  PERP_SEX + PERP_RACE,
  family=binomial,
  data=clean_incident_df,
  na.action = na.exclude)

print(levels(clean_incident_df$VIC_RACE))
```

```
[1] "WHITE"           "AM_INDIAN/ALASKAN" "ASIAN_PAC_ISLAND"
[4] "BLACK"           "BLACK HISPANIC"    "UNKNOWN"
[7] "WHITE HISPANIC"
```

```
print(levels(clean_incident_df$PERP_RACE))
```

```
[1] "WHITE"           "(null)"             "AM_INDIAN/ALASKAN"
[4] "ASIAN_PAC_ISLAND" "BLACK"              "BLACK HISPANIC"
[7] "UNKNOWN"         "WHITE HISPANIC"
```

```
print(levels(clean_incident_df$BORO))
```

```
[1] "BRONX"          "BROOKLYN"       "MANHATTAN"      "QUEENS"
[5] "STATEN ISLAND"
```

```
print(summary(model_vic_sex))
```

Call:

```
glm(formula = VIC_SEX ~ BORO + VIC_RACE + PERP_SEX + PERP_RACE,
     family = binomial, data = clean_incident_df, na.action = na.exclude)
```

Coefficients:

	Estimate	Std. Error	z value	Pr(> z)
(Intercept)	-1.06855	0.20611	-5.184	2.17e-07 ***
BOROBROOKLYN	0.16083	0.05205	3.090	0.002003 **
BOROMANHATTAN	0.11455	0.06716	1.706	0.088079 .
BOROQUEENS	0.17563	0.06462	2.718	0.006571 **
BOROSTATEN ISLAND	0.22120	0.11703	1.890	0.058742 .
VIC_RACEAM_INDIAN/ALASKAN	-0.60707	1.05697	-0.574	0.565729
VIC_RACEASIAN_PAC_ISLAND	-0.71881	0.19994	-3.595	0.000324 ***
VIC_RACEBLACK	-0.62712	0.11250	-5.574	2.49e-08 ***
VIC_RACEBLACK HISPANIC	-0.56914	0.12821	-4.439	9.04e-06 ***
VIC_RACEUNKNOWN	-0.20656	0.37435	-0.552	0.581100
VIC_RACEWHITE HISPANIC	-0.30462	0.11912	-2.557	0.010549 *
PERP_SEXM	-0.54537	0.12797	-4.262	2.03e-05 ***
PERP_SEXU	-0.48918	0.23420	-2.089	0.036736 *
PERP_RACE(null)	-0.11578	0.27836	-0.416	0.677451
PERP_RACEAM_INDIAN/ALASKAN	-8.69531	84.17442	-0.103	0.917724
PERP_RACEASIAN_PAC_ISLAND	0.20867	0.28233	0.739	0.459852
PERP_RACEBLACK	0.01052	0.17463	0.060	0.951985
PERP_RACEBLACK HISPANIC	-0.32055	0.19623	-1.634	0.102363
PERP_RACEUNKNOWN	-0.50673	0.25938	-1.954	0.050742 .
PERP_RACEWHITE HISPANIC	-0.25574	0.18346	-1.394	0.163326

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 18197 on 28561 degrees of freedom
Residual deviance: 17987 on 28542 degrees of freedom

AIC: 18027

Number of Fisher Scoring iterations: 9

Potential Bias

There are several sources of potential bias in this report:

- **Collection bias:** Some of the categorical variable options, such as the descriptions of locations or location categories, are limited, and some actual locations may be ambiguous. The racial categorization options are also somewhat limited compared to the true diversity of racial and ethnic backgrounds. The racial categorization is likely to have been chosen by someone other than the subject, so the selected race may not be accurate or reflect how the subject (victim or perpetrator) would self-identify.
- **Analysis bias:** There are likely many interesting insights in this data regarding race (relationship between race of perpetrator and victim, racial distribution of incidents throughout the boroughs, etc.), but I deliberately avoided any in-depth analysis of those categories. I believed it would be more complicated and nuanced than my current ability and available time would allow me to do well.

`sessionInfo()`

R version 4.4.1 (2024-06-14)
Platform: aarch64-apple-darwin20
Running under: macOS Sonoma 14.6.1

Matrix products: default
BLAS: /Library/Frameworks/R.framework/Versions/4.4-arm64/Resources/lib/libRblas.0.dylib
LAPACK: /Library/Frameworks/R.framework/Versions/4.4-arm64/Resources/lib/libRlapack.dylib;

locale:
[1] en_US.UTF-8/en_US.UTF-8/en_US.UTF-8/C/en_US.UTF-8/en_US.UTF-8

time zone: America/Denver
tzcode source: internal

attached base packages:
[1] stats graphics grDevices utils datasets methods base

other attached packages:

```
[1] broom_1.0.6      forecast_8.23.0 lubridate_1.9.3 forcats_1.0.0
[5] stringr_1.5.1    dplyr_1.1.4     purrr_1.0.2     readr_2.1.5
[9] tidyr_1.3.1      tibble_3.2.1    ggplot2_3.5.1   tidyverse_2.0.0
```

loaded via a namespace (and not attached):

```
[1] utf8_1.2.4      generics_0.1.3  stringi_1.8.4   lattice_0.22-6
[5] hms_1.1.3       digest_0.6.36   magrittr_2.0.3  evaluate_0.24.0
[9] grid_4.4.1      timechange_0.3.0 fastmap_1.2.0   jsonlite_1.8.8
[13] backports_1.5.0 nnet_7.3-19     fansi_1.0.6     scales_1.3.0
[17] cli_3.6.3       rlang_1.1.4     munsell_0.5.1   withr_3.0.1
[21] yaml_2.3.10     tools_4.4.1     parallel_4.4.1  tzdb_0.4.0
[25] colorspace_2.1-1 curl_5.2.2      vctrs_0.6.5     R6_2.5.1
[29] zoo_1.8-12      lifecycle_1.0.4 tseries_0.10-58 urca_1.3-4
[33] pkgconfig_2.0.3 pillar_1.9.0    gtable_0.3.5    quantmod_0.4.26
[37] glue_1.7.0      Rcpp_1.0.13     xfun_0.46       lmtest_0.9-40
[41] tidyselect_1.2.1 rstudioapi_0.16.0 knitr_1.48      farver_2.1.2
[45] nlme_3.1-164    htmltools_0.5.8.1 labeling_0.4.3  xts_0.14.0
[49] rmarkdown_2.28  timeDate_4041.110 fracdiff_1.5-3  compiler_4.4.1
[53] quadprog_1.5-8  TTR_0.24.4
```